



Formation

# Introduction au Deep Learning

Séquence 8

“Données creuses et embedding, données séquentielles et réseaux de Neurones Récurrents (RNN)”



FIDLE





Formation

# Introduction au Deep Learning

<https://fidle.cnrs.fr>

Powered by CNRS CRIC,  
and UGA DGDSI of Grenoble, thanks !



Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0)  
<https://creativecommons.org/licenses/by-nc-nd/4.0/>

-  Course materials (pdf)
-  Practical work environment
-  Corrected notebooks
-  Videos (YouTube)



(\*) Procedure via Docket or pip  
Remember to get the latest version !



Formation  
**Introduction au**  
**Deep Learning**

Questions and answers :

<https://fidle.cnrs.fr/q2a>



**Accompanied by :**  
IA Support (dream) Team of IDRIS

**Directed by :**  
Agathe, Baptiste et Yanis - UGA/DAPI  
Thibaut, Kamel - IDRIS



Formation

# Introduction au Deep Learning



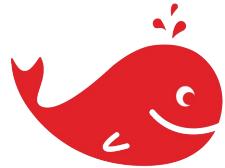
<https://fidle.cnrs.fr/listeinfo>  
Fidle information list

**Agoria**

<http://fidle.cnrs.fr/agoria>  
AI exchange list



[agoria@grenoble.cnrs.fr](mailto:agoria@grenoble.cnrs.fr)



FIDLE

Formation

# Introduction au Deep Learning

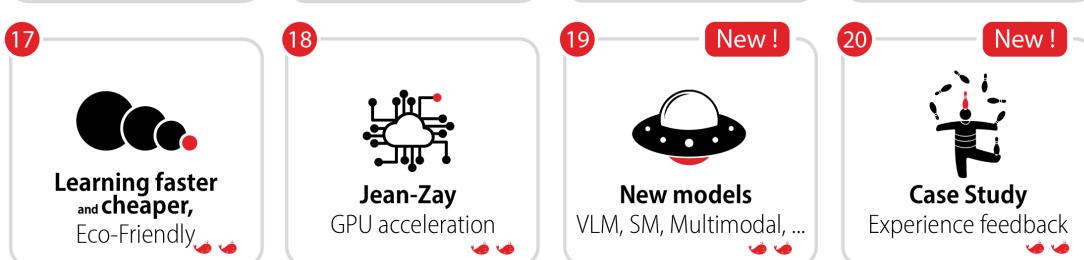
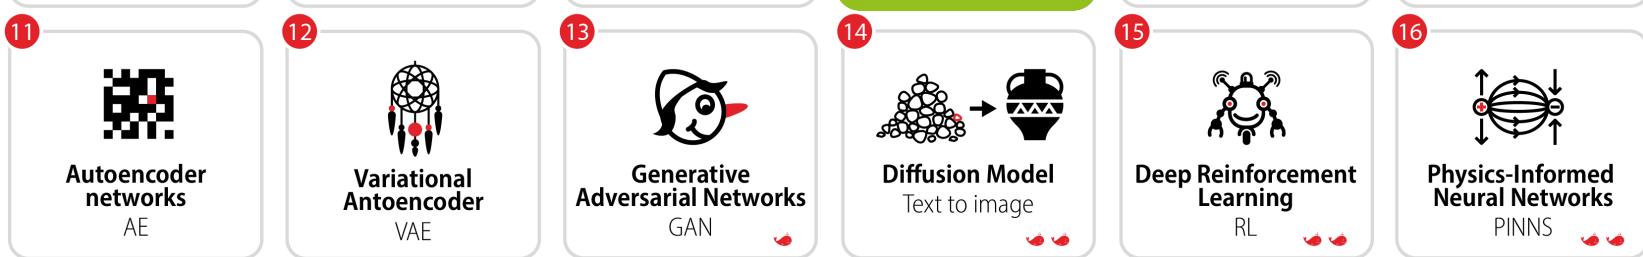
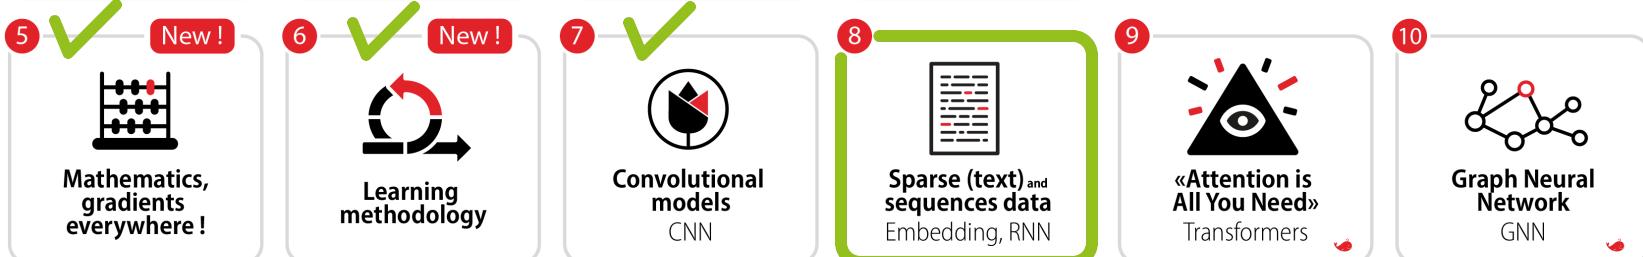
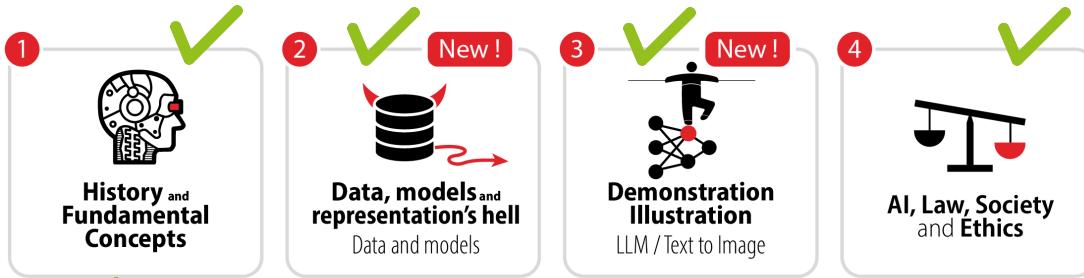


GROUPECALCUL

<https://listes.services.cnrs.fr/wws/info/devlog>  
List of ESR\* « Software developers » group

<https://listes.math.cnrs.fr/wws/info/calcul>  
List of ESR\* « Calcul » group

## Bases, Concepts et Enjeux



## L'IA comme un outil,

## Acteur de l'IA



8



**Sparse (text) and sequences data**  
Embedding, RNN

**Part 1**



**Sparse data (text)**  
Embedding

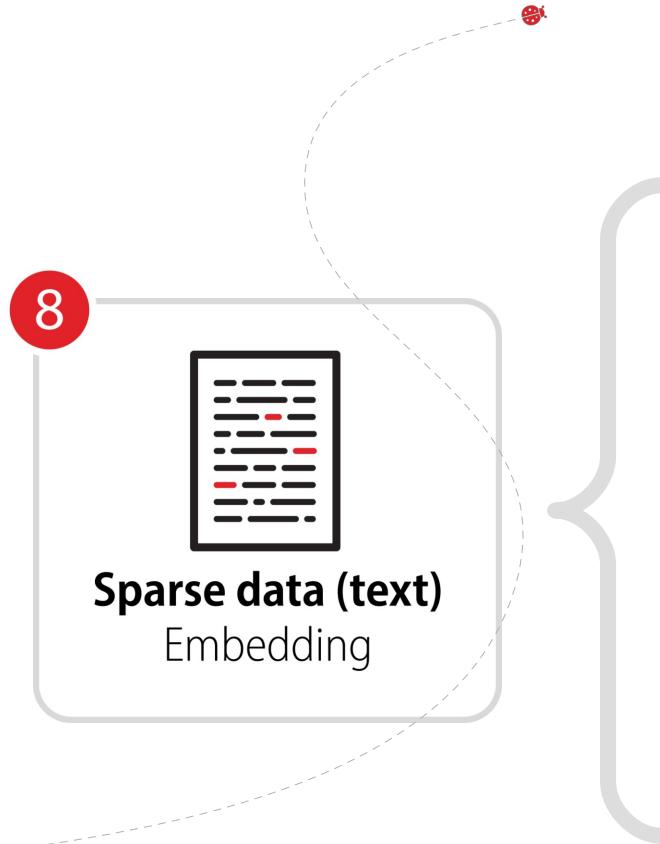
**Part 2**



9



**«Attention is All You Need»**  
Transformers



## Part 1

- 1 **Text encoding**
  - From text to tensor
  - One hot encoding
  - Embedding
- 2 **Example 1 : IMDB1**
  - Sentiment analysis with **one-hot encoding**
- 3 **Example 2 : IMDB2/3/4**
  - Sentiment analysis with **embedding**

8

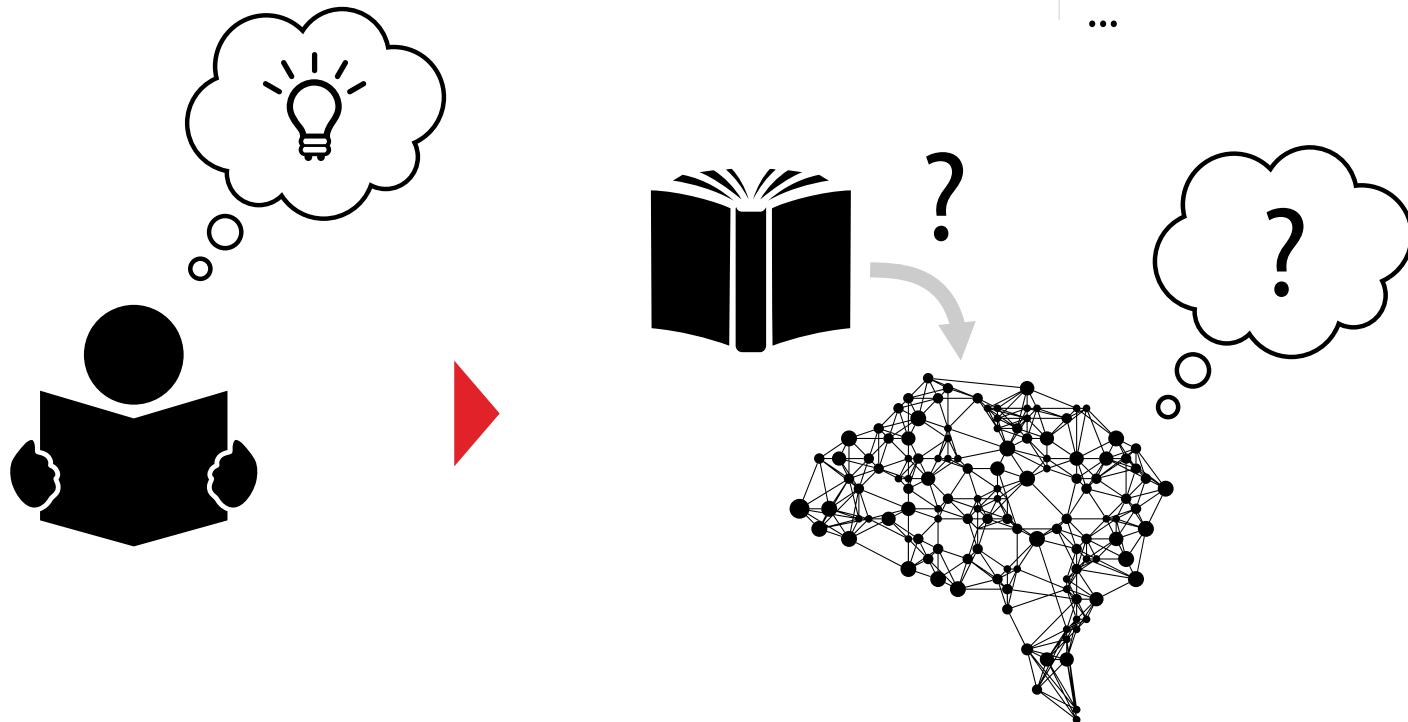


**Sparse (text) and  
sequences data**  
Embedding, RNN

## Part 1

- 1 **Text encoding**
  - From text to tensor
  - One hot encoding
  - Embedding
- 2 **Example 1 : IMDB1**
  - Sentiment analysis with **one-hot encoding**
- 3 **Example 2 : IMDB2/3/4**
  - Sentiment analysis with **embedding**

# How to feed a neural network with text ?



Variable size  
Large size space  
Conceptual contents ;-)  
...

« I've never seen a movie like this before. »



How to build a descriptor for this kind of data (text, DNA, ...)?

# One hot encoding

« I've never seen a movie like this before. »



**Tokenization**  
(words)

Dictionary

|   |           |
|---|-----------|
| 0 | a         |
| 1 | before    |
| 2 | fantastic |
| 3 | i've      |
| 4 | is        |
| 5 | like      |
| 6 | movie     |
| 7 | never     |
| 8 | seen      |
| 9 | this      |

["i've", 'never', 'seen', 'a', 'movie', 'like', 'this', 'before']



The word order in the dictionary may have some meaning. For example according to their rate of use.

# One hot encoding

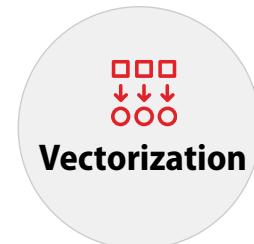
« I've never seen a movie like this before. »

Dictionary

|   |           |
|---|-----------|
| 0 | a         |
| 1 | before    |
| 2 | fantastic |
| 3 | i've      |
| 4 | is        |
| 5 | like      |
| 6 | movie     |
| 7 | never     |
| 8 | seen      |
| 9 | this      |



["i've", 'never', 'seen', 'a', 'movie', 'like', 'this', 'before']  
[ 3, 7, 8, 0, 6, 5, 9, 1 ]



! The values associated with each word are **meaningless**: words with a contiguous subscript are typically unrelated.

# One hot encoding

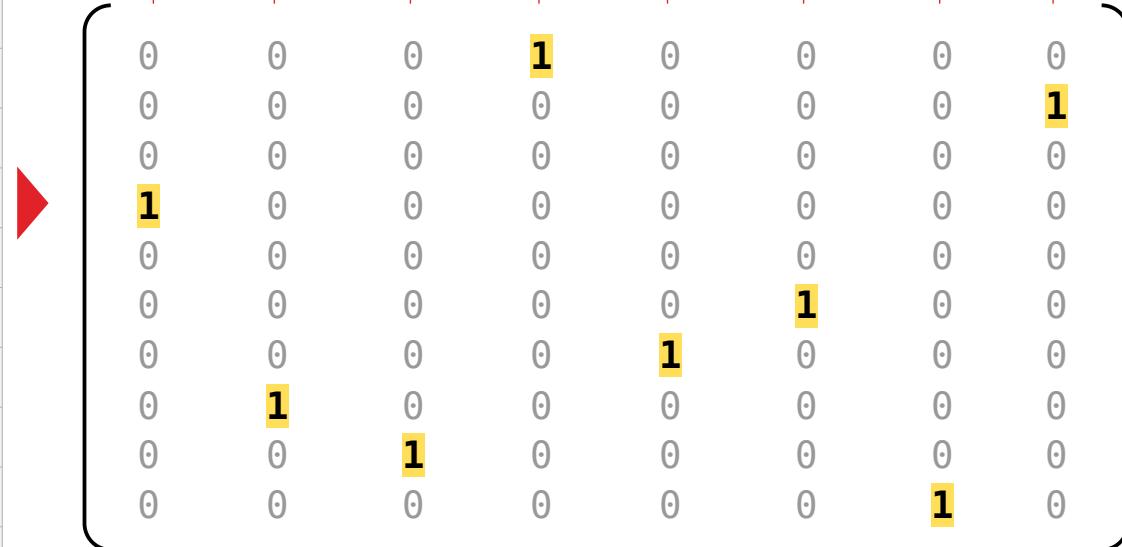
« I've never seen a movie like this before. »

Dictionary

|   |           |
|---|-----------|
| 0 | a         |
| 1 | before    |
| 2 | fantastic |
| 3 | i've      |
| 4 | is        |
| 5 | like      |
| 6 | movie     |
| 7 | never     |
| 8 | seen      |
| 9 | this      |

["i've", 'never', 'seen', 'a', 'movie', 'like', 'this', 'before']

[ 3, 7, 8, 0, 6, 5, 9, 1 ]



Tokenization  
(words)



Vectorization

001000

One hot  
encoding

# One hot encoding

« I've never seen a movie like this before. »



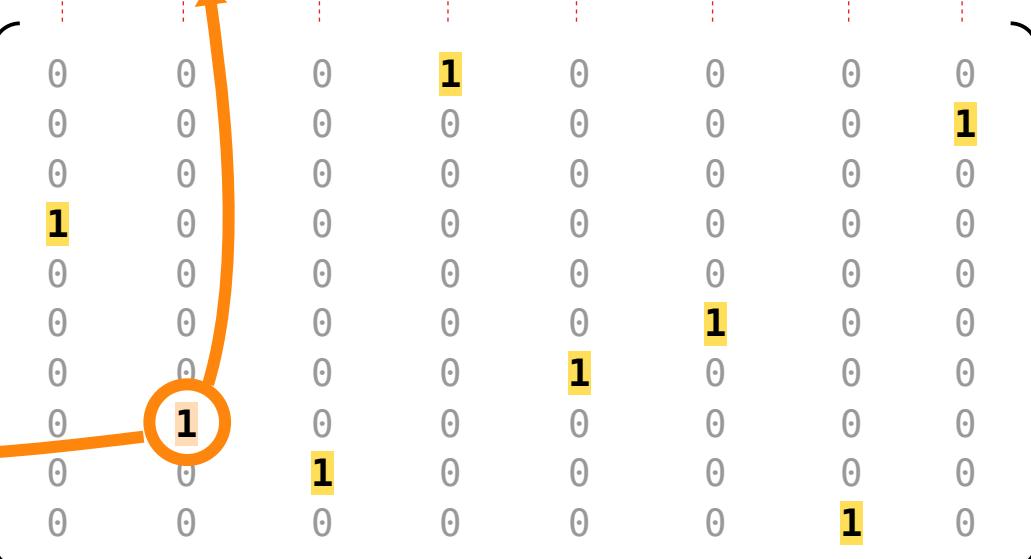
Tokenization  
(words)

Dictionary

|   |           |
|---|-----------|
| 0 | a         |
| 1 | before    |
| 2 | fantastic |
| 3 | i've      |
| 4 | is        |
| 5 | like      |
| 6 | movie     |
| 7 | never     |
| 8 | seen      |
| 9 | this      |

["i've", 'never', 'seen', 'a', 'movie', 'like', 'this', 'before']

[ 3, 7, 8, 0, 6, 5, 9, 1 ]



Vectorization

001000

One hot  
encoding



Each vector is independent.  
Each word has its own dimension.

# One hot encoding



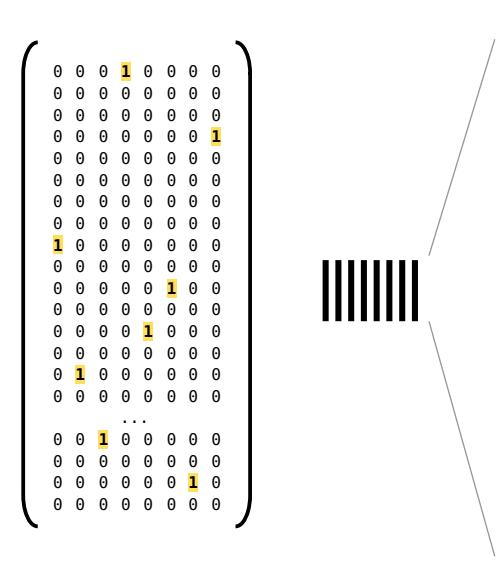
# Limits of full one-hot encoding : Size !

## Example :

Dictionary = 80 000 words

**Sentence = 300 words**

2 answers : { One hot vector  
Embedding

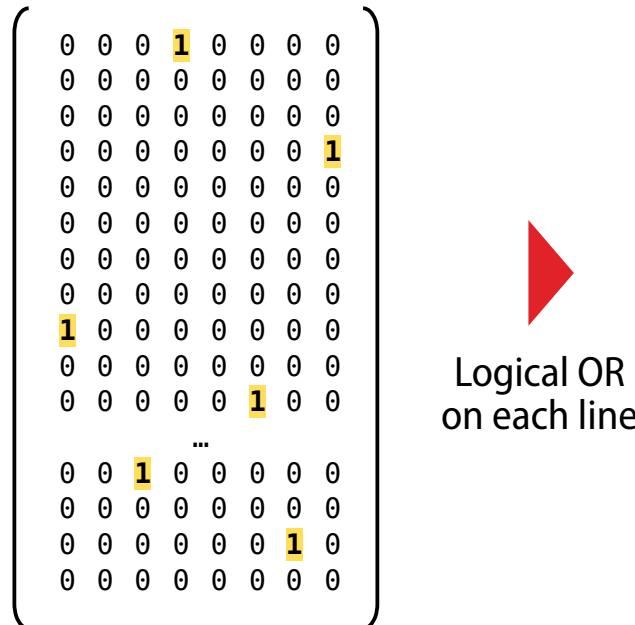


Need :  
 **$24 \cdot 10^6$**   
Parameters !

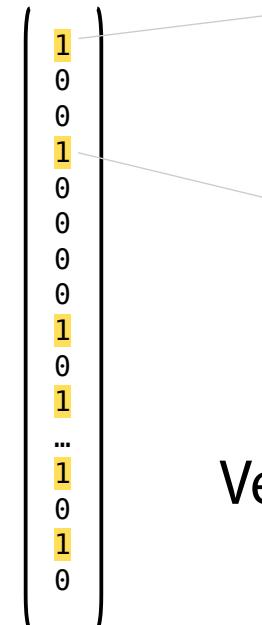


# One hot vector encoding

## Solution 1 : one hot vector



Logical OR  
on each line



- 1 This word is **present** in the sentence
  - 0 This word is **not present** in the sentence

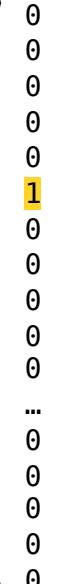
# Vector size = Dictionary size



The size no longer depends on the length of the sentence.  
The size is drastically reduced, but we lose the words order.



## Solution 2 : Using embedding

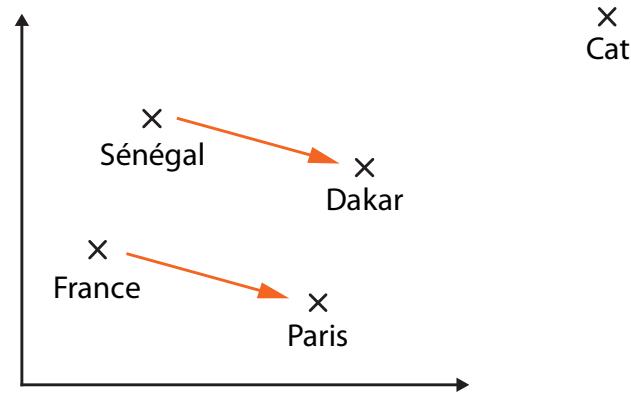
(Word<sub>i</sub>) 

▶  $(W_i) \begin{pmatrix} 0.7 \\ 4.2 \\ 1.6 \end{pmatrix}$

Short dense vector,  
whose value carries a  
semantic meaning of  
the word.

Long sparse vector,  
whose value just indicates  
membership in the dictionary

Semantic meaning, allows to perform calculations, such as :



$\text{dist}(\text{'France'}, \text{'Cat'}) >> \text{dist}(\text{'France'}, \text{'Sénégal'})$

$\text{vect}(\text{'France'}, \text{'Paris'}) \approx \text{vect}(\text{'Sénégal'}, \text{'Dakar'})$

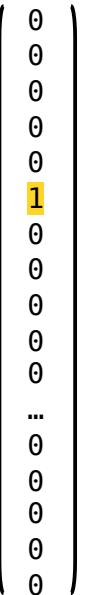
$\text{'Paris'} = \text{'France'} + \text{vect}(\text{'Sénégal'}, \text{'Dakar'})$

« Paris est à la France ce que Dakar est au Sénégal »



## Solution 2 : Using embedding

(Word<sub>i</sub>)



$$(W_i) \begin{pmatrix} 0.7 \\ 4.2 \\ 1.6 \end{pmatrix}$$

Short dense vector,  
whose value carries a  
semantic meaning of  
the word.

Long sparse vector,  
whose value just indicates  
membership in the dictionary

Number of parameters required :

Example :

Embedding size :

200

Sentence size :

300 words



Need only :

60.000

Parameters :-)



Embedding techniques :

Contextual  
embedding

Global  
embedding

Keras embedding

CBOW,  
SG,  
GloVe,  
etc.

## Embedding layers in Keras

Usable as a **simple layer**

This layer will constitute a **dictionary of vectors** which will be **optimized during the learning process**, according to the expected result and not to the pure semantics.

Keras embedding is therefore adapted to **classification**, but will not be able to take into consideration, for example, semantic similarities (such as identifying 2 sentences with the same semantic meaning).

The **output** of the layer is a **set of vectors**.

## Embedding layers in Keras

(french)

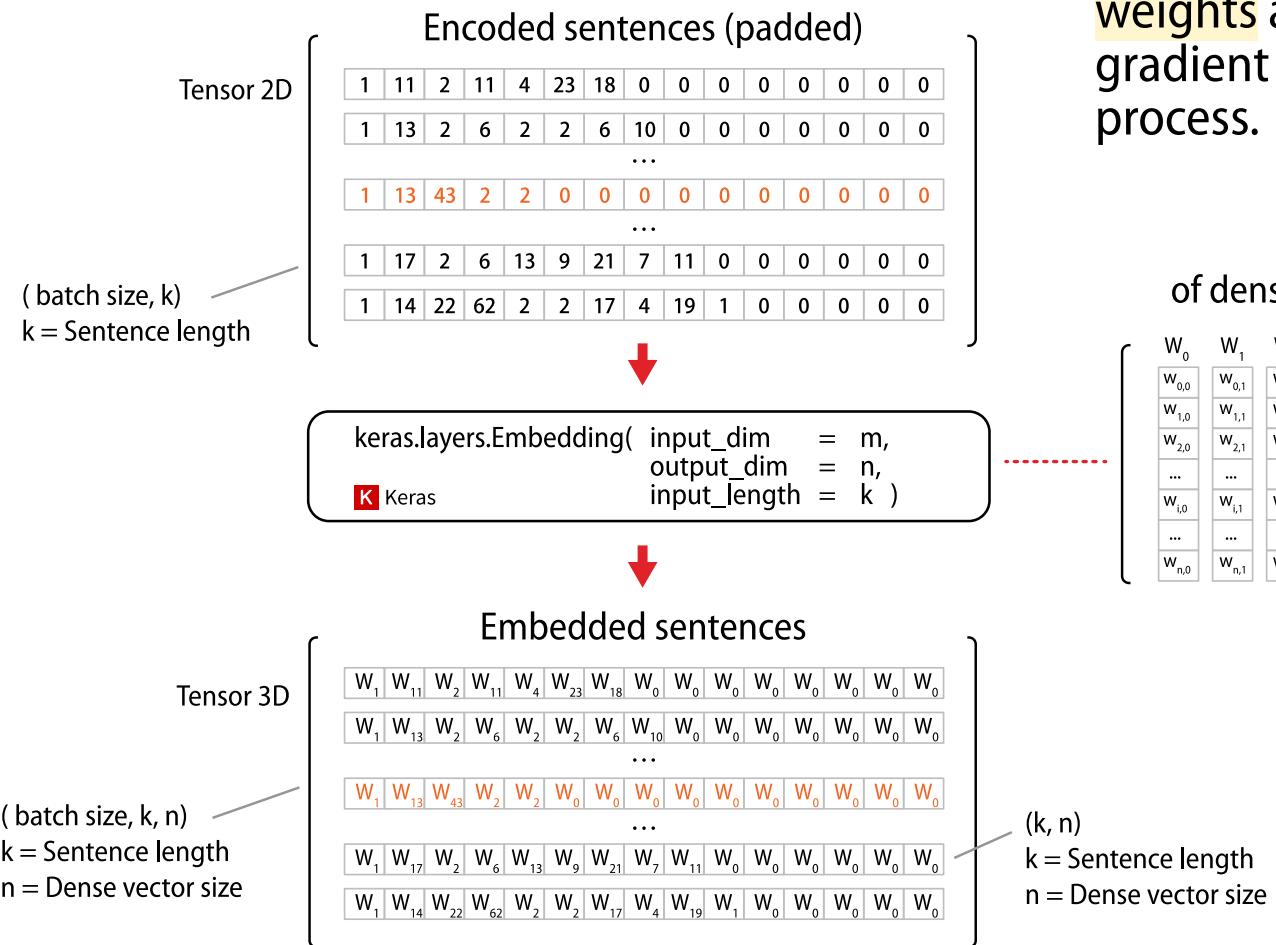
Utilisable comme une **simple couche**

Cette couche va constituer un **dictionnaire de vecteurs** qu'elle optimisera au cours de l'apprentissage, en **fonction du résultat attendu** et non de la sémantique pure.

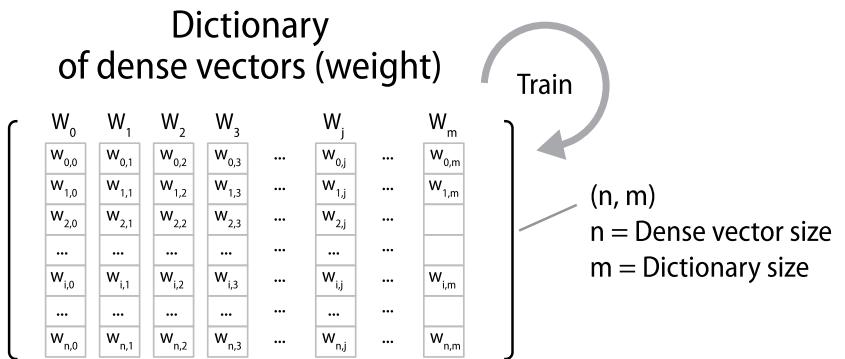
L'embedding Keras est donc adapté à la **classification**, mais ne rendra pas compte, par exemple, de similarités sémantiques (comme identifier 2 phrases ayant un même sens sémantique).

La **sortie** de la couche est un **ensemble de vecteurs**

# Keras embedding layer



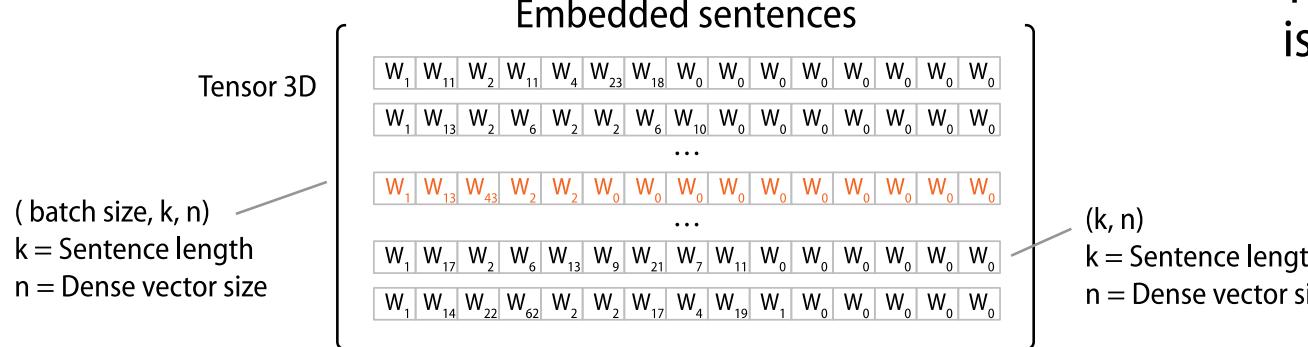
Dense vector components are managed as weights and updated through the standard gradient descent and backpropagation process.



# Keras embedding layer



The order of the words  
is therefore ignored !

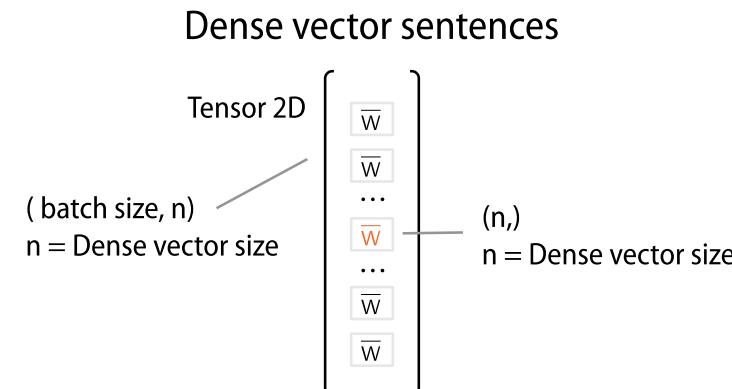


-----  $\overline{W^{(s)}} = \frac{1}{k} \sum_{i=1}^k W_i^{(s)}$

Where :

$k$  is sentence length  
 $W_i^{(s)}$  is word vector i of sentence (s)

Each sentence is  
here encoded by a  
unique dense  
vector which is the  
average of the  
words composing  
the sentence.



## Word2Vec

This approach aims to build **dictionaries** whose vector representation of words is based on **context** and therefore on **semantics**.

Dictionaries built from large corpora are available.

Two models:

- **Continuous Bag-of-Words** (CBOW),
- **Skip-Gram** (SG).

These approaches are historically interesting, but now a bit outdated... ;-)

### Efficient Estimation of Word Representations in Vector Space

Tomas Mikolov, Kai Chen, Greg Corrado, Jeffrey Dean, 2013

<https://doi.org/10.48550/arXiv.1301.3781>

CBOW : Continuous Bag of Words - Embedding based on the prediction of the word according to its context.

SG : Skip-gram - Embedding based on context prediction from the word.



## Word2Vec

Approche ayant pour objectif de constituer des **dictionnaires** dont la représentation vectorielle des mots est basée sur le **contexte** et donc de la **sémantique**.

Des dictionnaires construits à partir de gros corpus sont disponibles.

Deux modèles :

- **Continuous Bag-of-Words** (CBOW),
- **Skip-Gram** (SG).

Ces approches sont historiquement intéressantes, mais désormais un peu dépassées... ;-)

### Efficient Estimation of Word Representations in Vector Space

Tomas Mikolov, Kai Chen, Greg Corrado, Jeffrey Dean, 2013

<https://doi.org/10.48550/arXiv.1301.3781>

CBOW : Continuous Bag of Words - Embedding based on the prediction of the word according to its context.

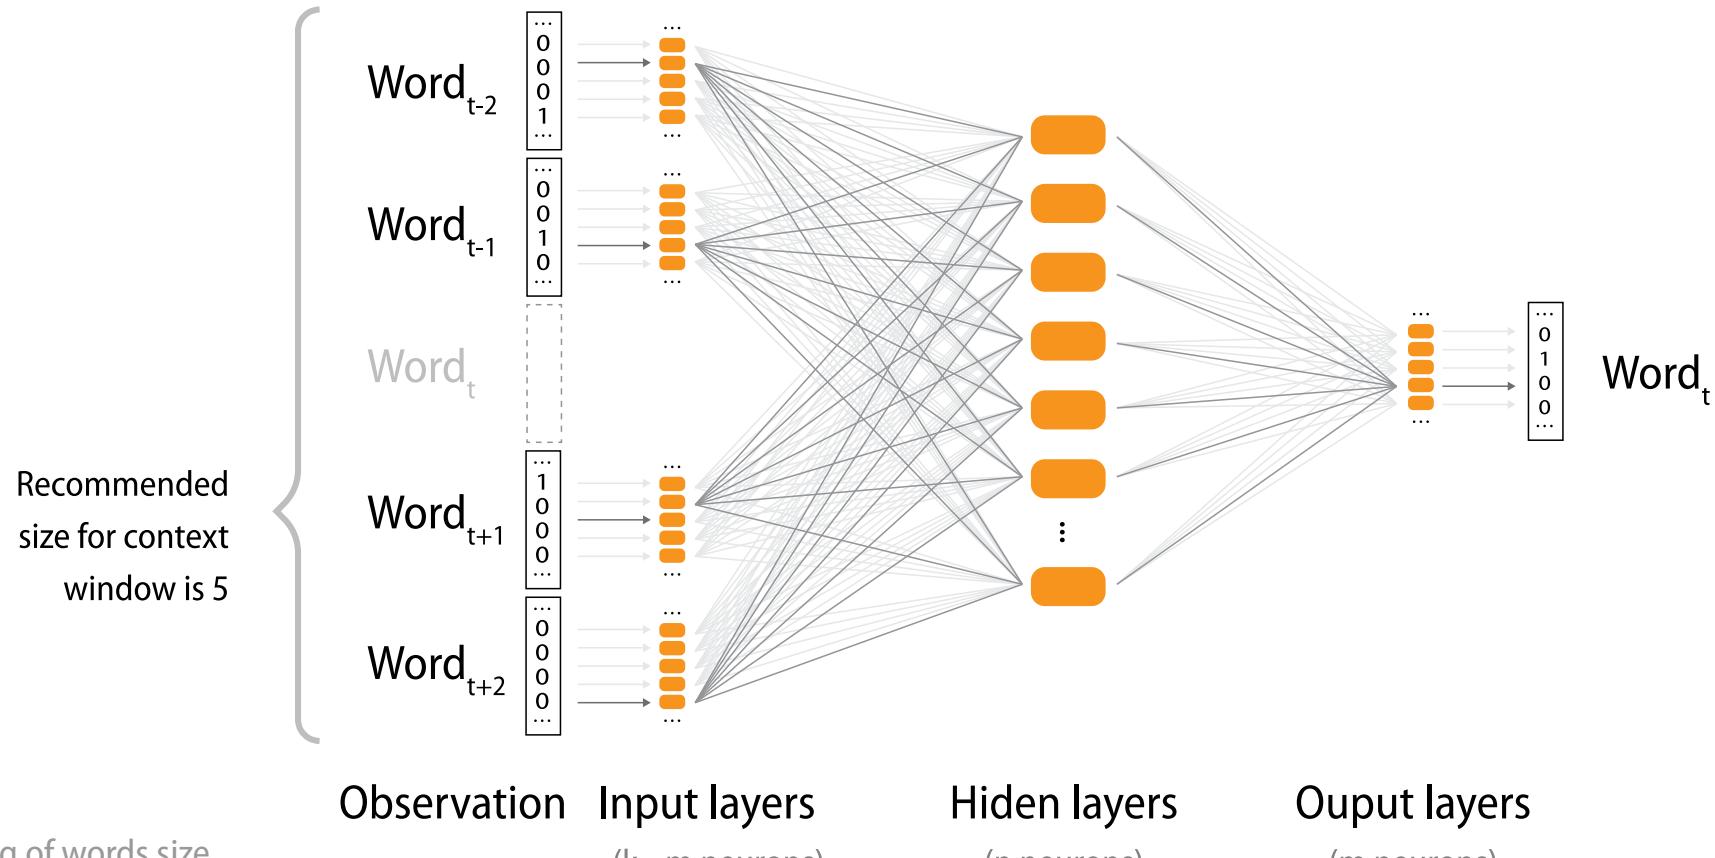
SG : Skip-gram - Embedding based on context prediction from the word.



# Continuous Bag-of-Words



The objective of training is to  
(re)find a word from its context.

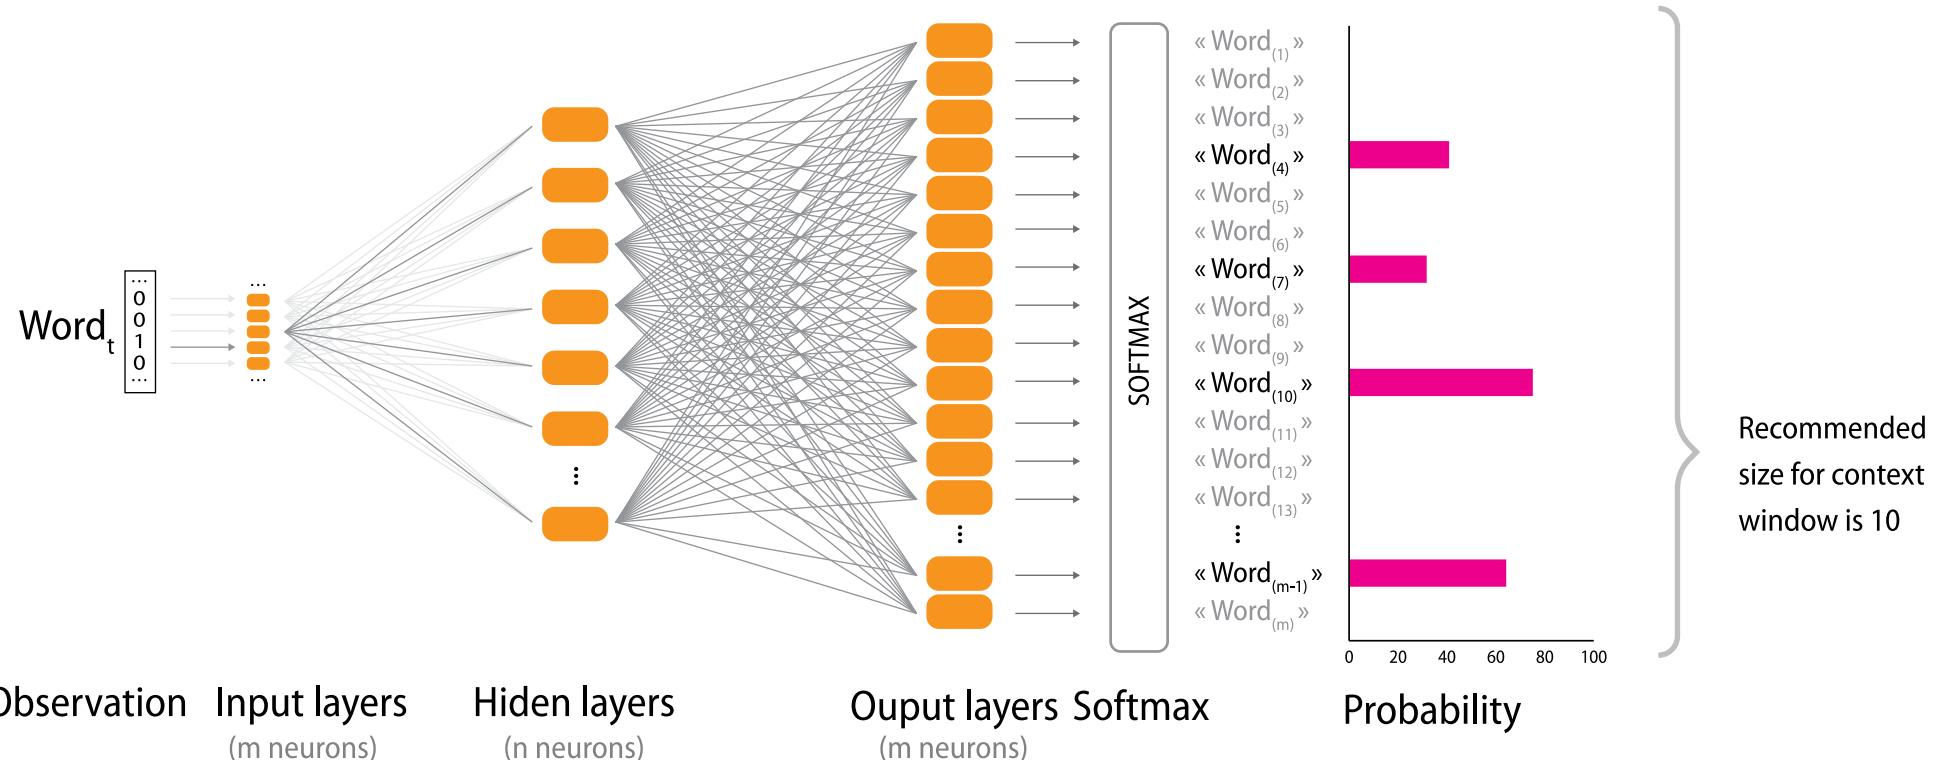


$k$  = Bag of words size  
 $n$  = Dense vector size  
 $m$  = Dictionary size

# Skip-Gram



The objective of training is to  
(re)find the context from the word.



Observation    Input layers  
( $m$  neurons)

Hidden layers  
( $n$  neurons)

Output layers    Softmax  
( $m$  neurons)

Probability

$k$  = Bag of words size  
 $n$  = Dense vector size  
 $m$  = Dictionary size

## (Flau)BERT

2018

BERT est une représentation du langage proposée par Google en 2018, permettant de prendre en compte la dimension contextuelle du langage (« Avocat », peut être un fruit ou un juriste..).

FlauBERT<sup>2</sup> est une adaptation de l'algorithme au français.

**BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding.**

Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova

<https://arxiv.org/abs/1810.04805>

To be continued !

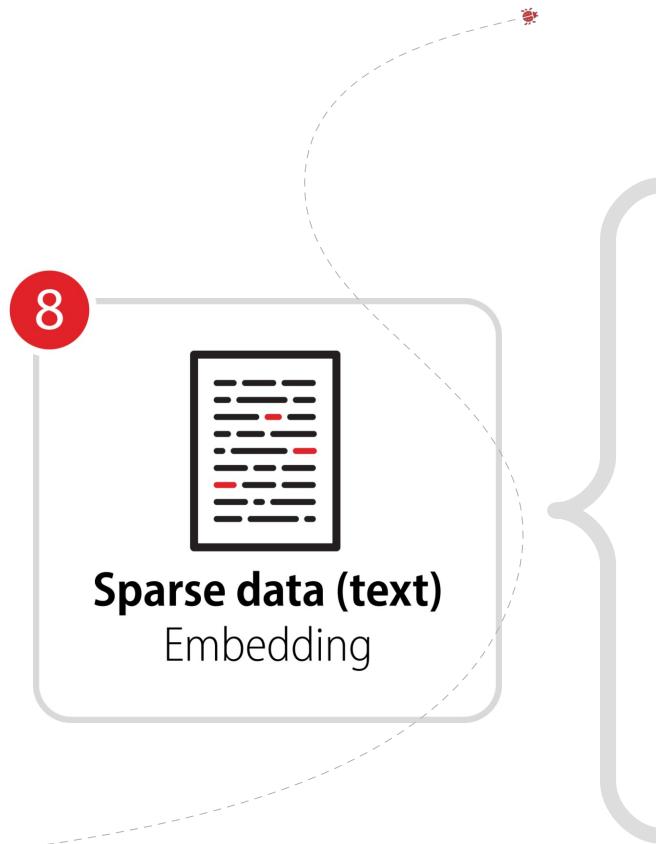
**FlauBERT: Unsupervised Language Model Pre-training for French.**

Hang Le, Loïc Vial, Jibril Frej, Vincent Segonne, Maximin Coavoux,

Benjamin Lecouteux, Alexandre Allauzen, Benoît Crabbé,

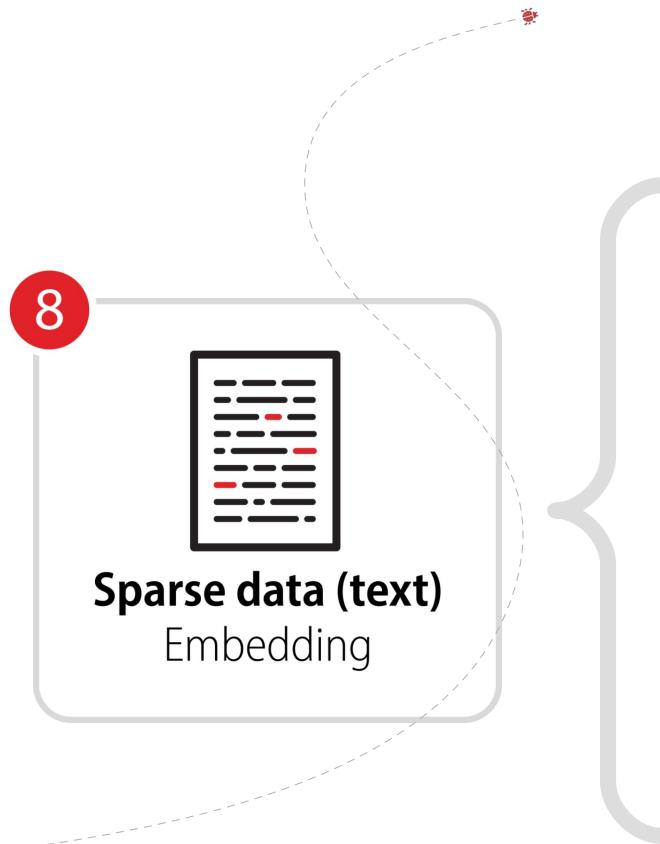
Laurent Besacier, Didier Schwab

<https://arxiv.org/abs/1912.05372>



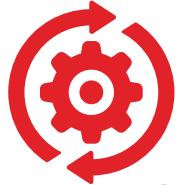
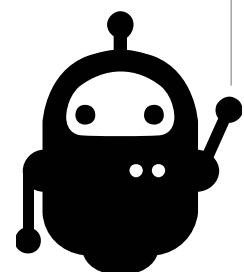
## Part 1

- 1 **Text encoding**
  - From text to tensor
  - One hot encoding
  - Embedding
- 2 **Example 1 : IMDB1**
  - Sentiment analysis with **one-hot encoding**
- 3 **Example 2 : IMDB2/3/4**
  - Sentiment analysis with **embedding**



## Part 1

- 1 **Text encoding**
  - From text to tensor
  - One hot encoding
  - Embedding
- 2 **Example 1 : IMDB1**
  - Sentiment analysis with **one-hot encoding**
- 3 **Example 2 : IMDB2/3/4**
  - Sentiment analysis with **embedding**



# One-Hot encoding with IMDB

Notebook : [K3IMDB1]

## **Objective :**

Guess whether a film review is positive or not based on the analysis of the text, using One-Hot encoding.

## **Dataset :**

The IMDB dataset is composed of 50,000 film reviews from the site of the same name.

<https://www.imdb.com/>



Import  
and init

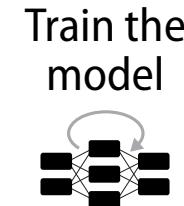
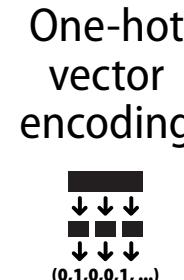
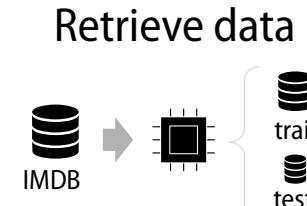
START

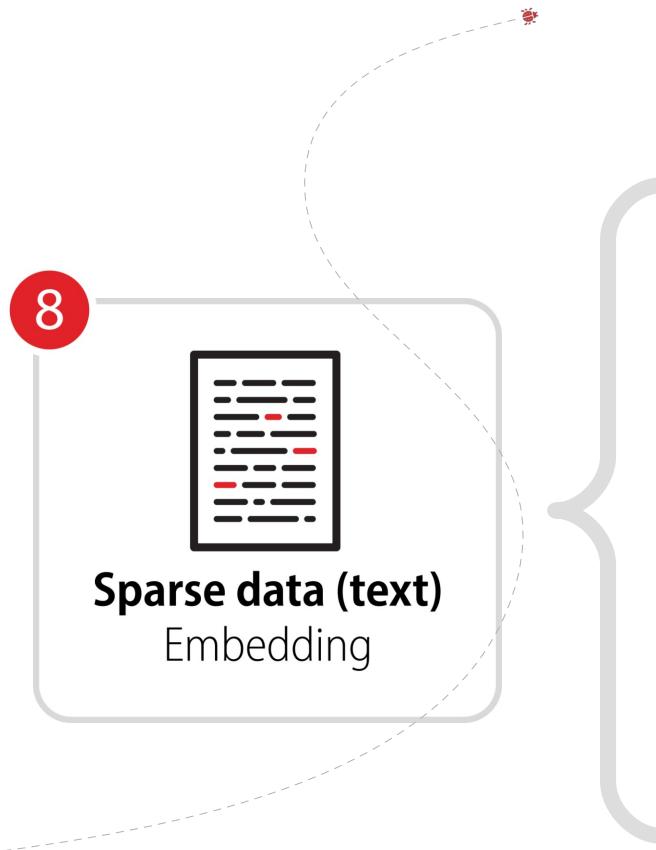
Understanding  
hot-one  
encodingAbout our  
datasetBuild the  
model

Evaluate



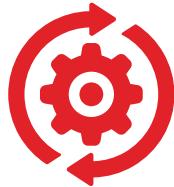
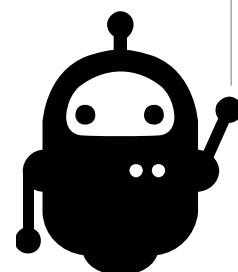
**Objectives :**  
Analysis of  
film reviews  
with  
**One-Hot**  
encoding





## Part 1

- 1 **Text encoding**
  - From text to tensor
  - One hot encoding
  - Embedding
- 2 **Example 1 : IMDB1**
  - Sentiment analysis with **one-hot encoding**
- 3 **Example 2 : IMDB2/3/4**
  - Sentiment analysis with **embedding**



# Text embedding with IMDB

Notebook : [K3IMDB2-4]

## **Objective :**

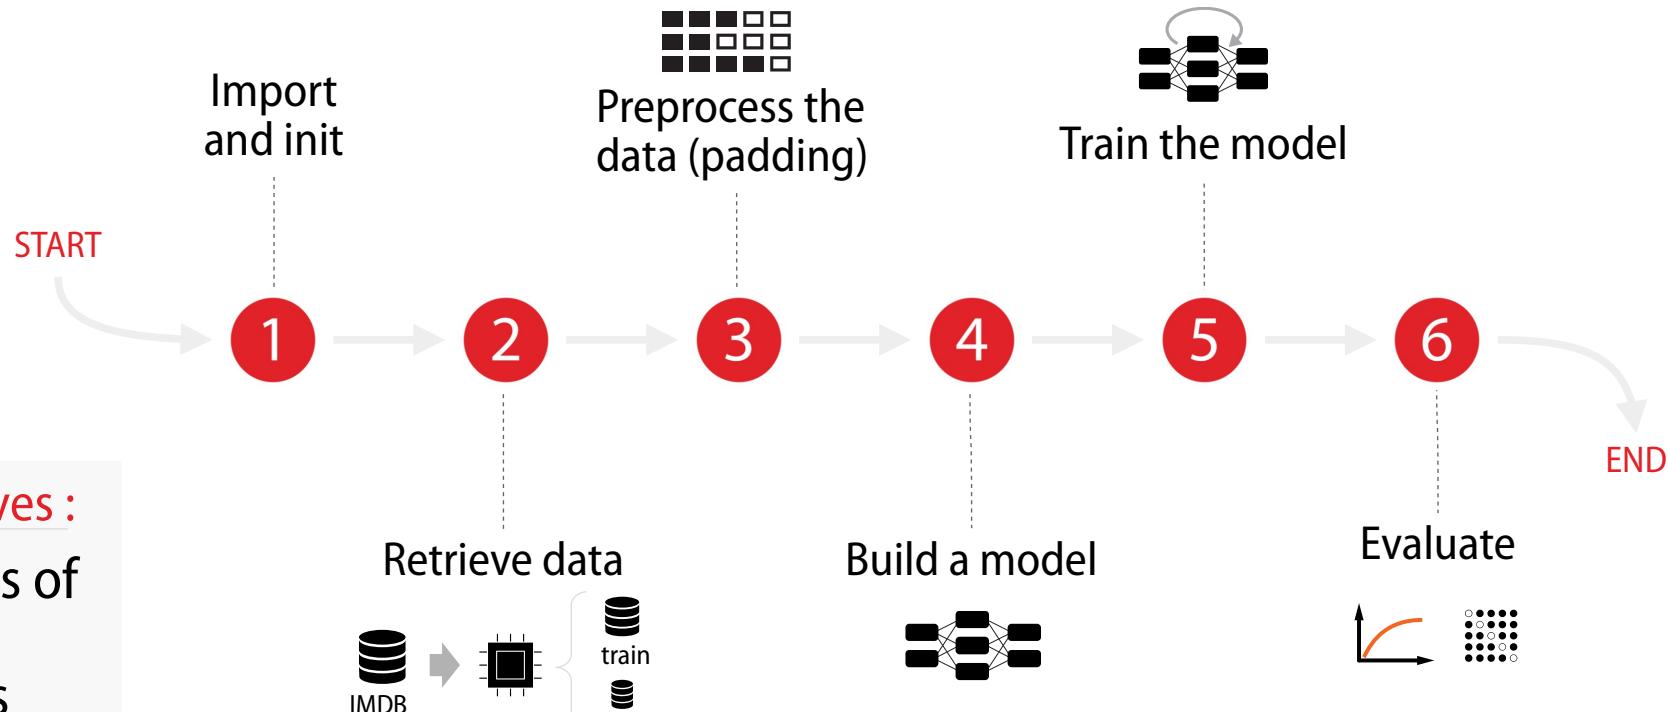
Guess whether a film review is positive or not based on the analysis of the text, using embedding.

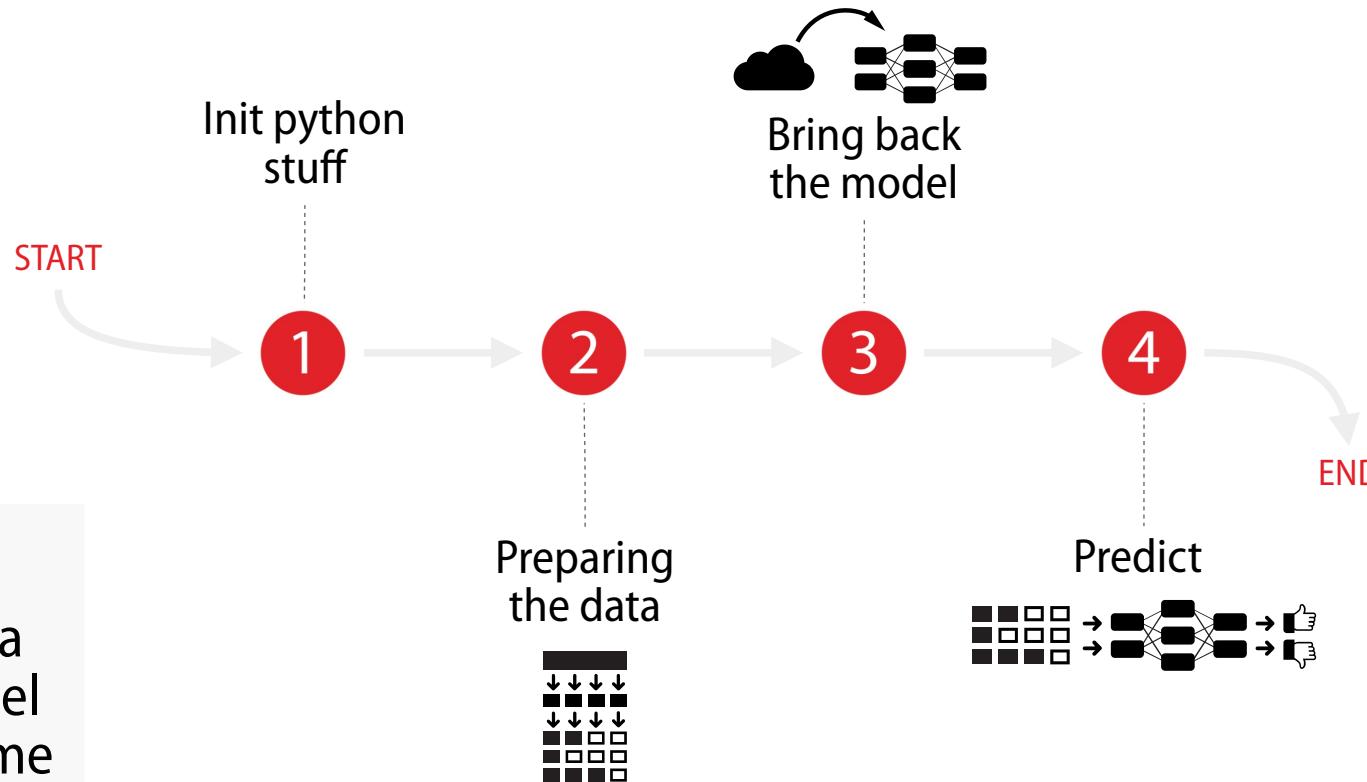
## **Dataset :**

The IMDB dataset is composed of 50,000 film reviews from the site of the same name.

<https://www.imdb.com/>



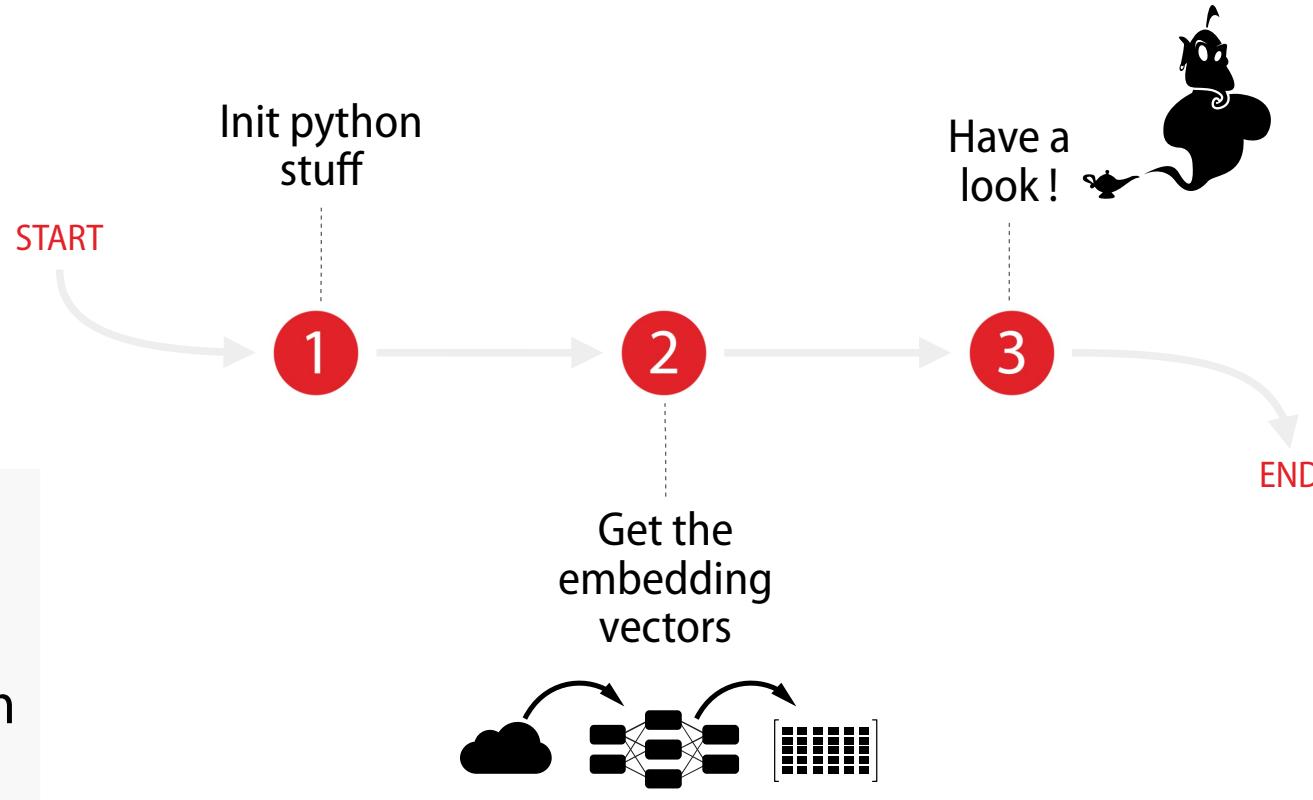




### Objectives :

Retrieving a saved model to treat some new reviews





8



**Sparse (text) and sequences data**  
Embedding, RNN

**Part 1**



**Sparse data (text)**  
Embedding

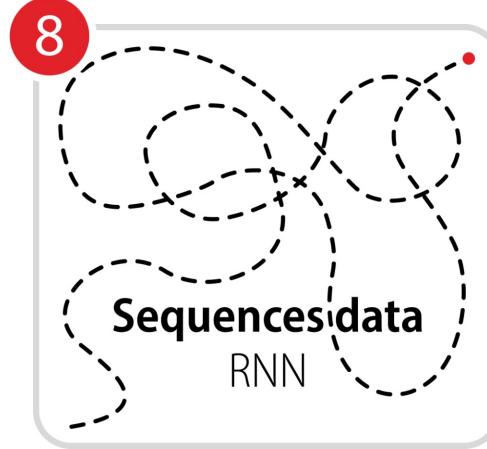
**Part 2**



9



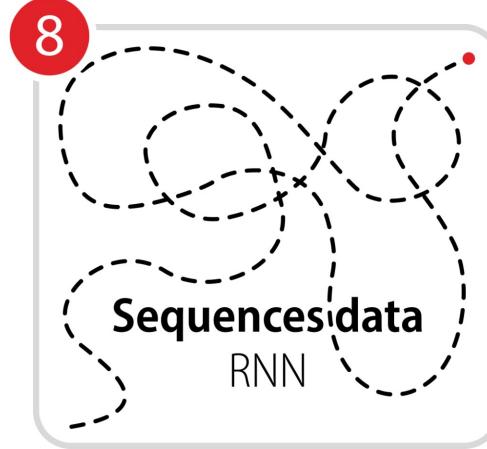
**«Attention is All You Need»**  
Transformers



## Part 2

- 1 **Sequences data**
  - Recurrent Neural Network
  - LSTM and GRU
- 2 **Example 1 : LadyBug**
  - Prediction of a virtual trajectory





## Part 2

1

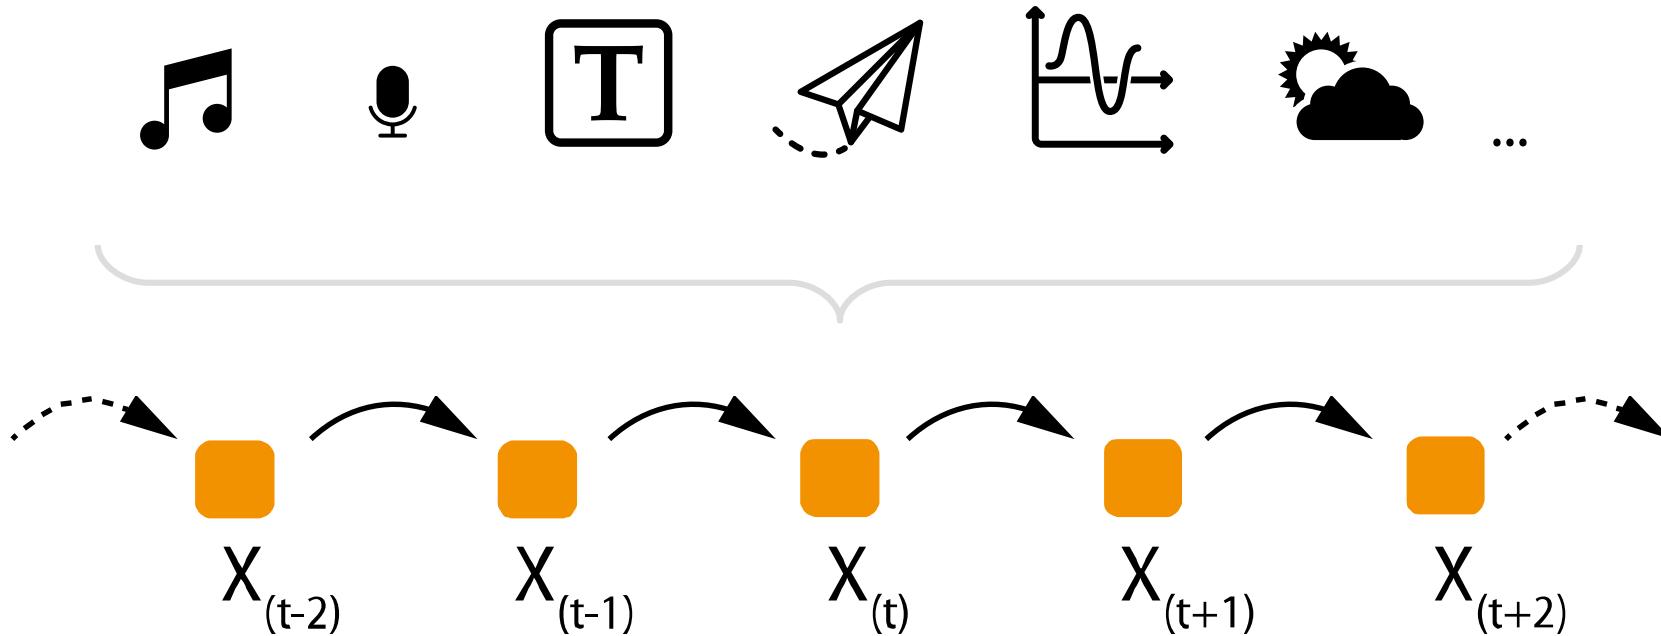
**Sequences data**  
→ Recurrent Neural Network  
→ LSTM and GRU

2

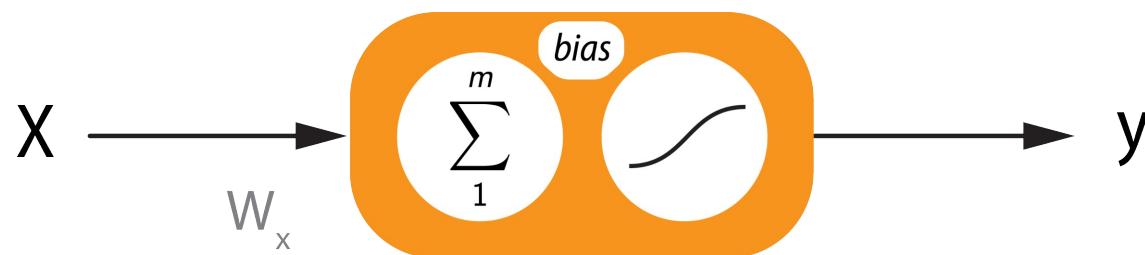
**Example 1 : LadyBug**  
→ Prediction of a virtual trajectory



# What if the world was just one big sequence ?

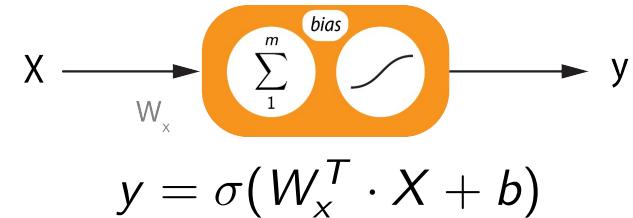
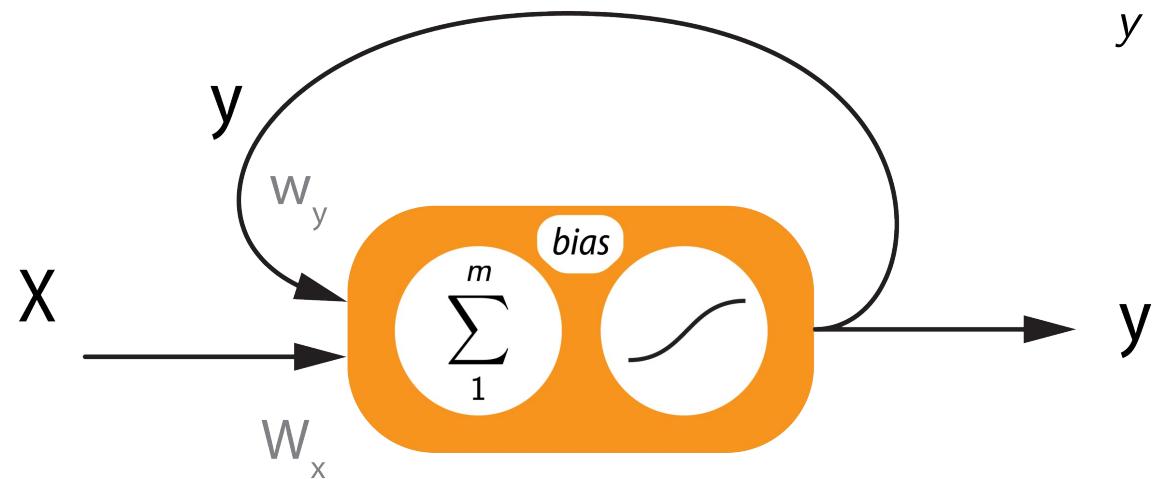


## Classical neuron.



$$y = \sigma(W_x^T \cdot X + b)$$

## Recurrent neuron.



$$y_{(t)} = \sigma(W_x^T \cdot X_{(t)} + w_y \cdot y_{(t-1)} + b)$$

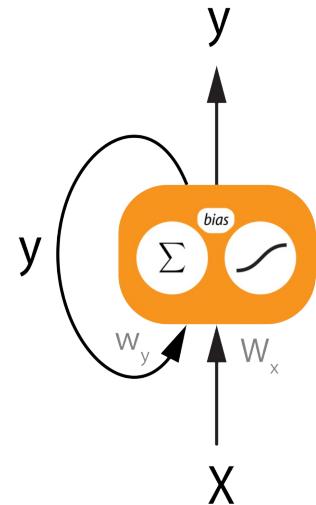
**w<sub>y</sub>** : is a scalar

**W<sub>x</sub>** : is a vector

**b** : is a scalar

**y** : is a scalar

## Recurrent neuron.



$$y(t) = \sigma(W_x^T \cdot X_{(t)} + w_y \cdot y_{(t-1)} + b)$$

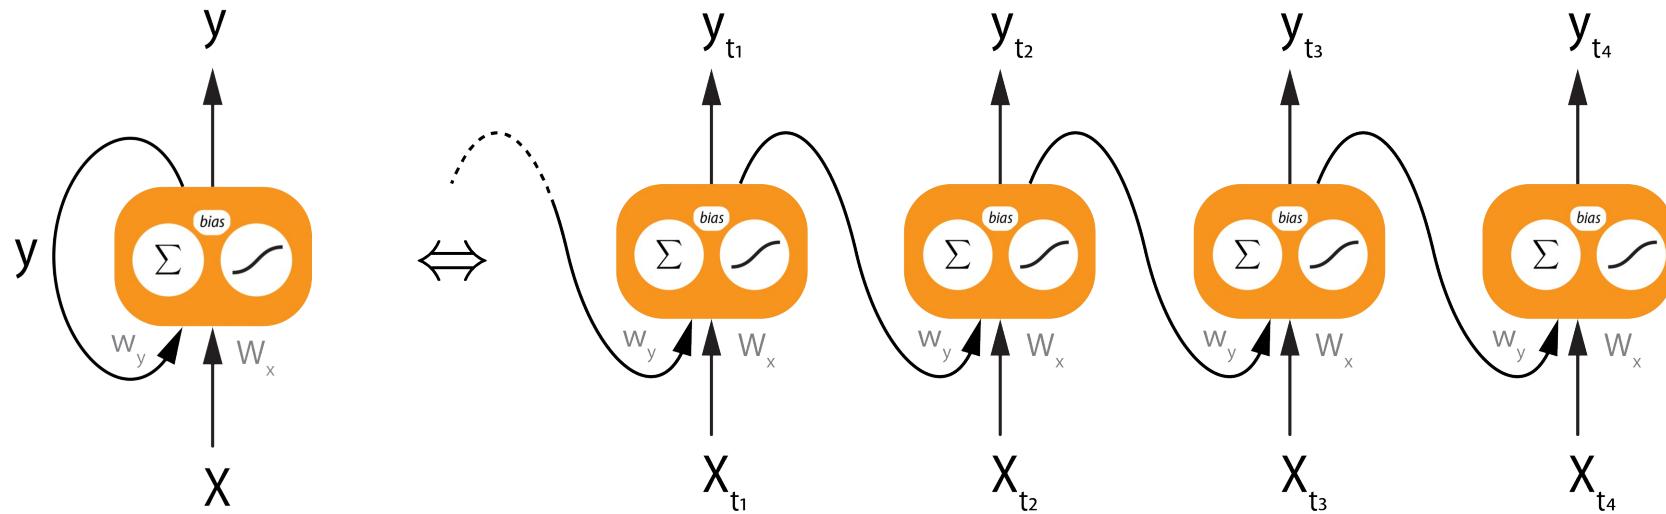
**w<sub>y</sub>** : is a scalar

**W<sub>x</sub>** : is a vector

**b** : is a scalar

**y** : is a scalar

## Recurrent neuron.

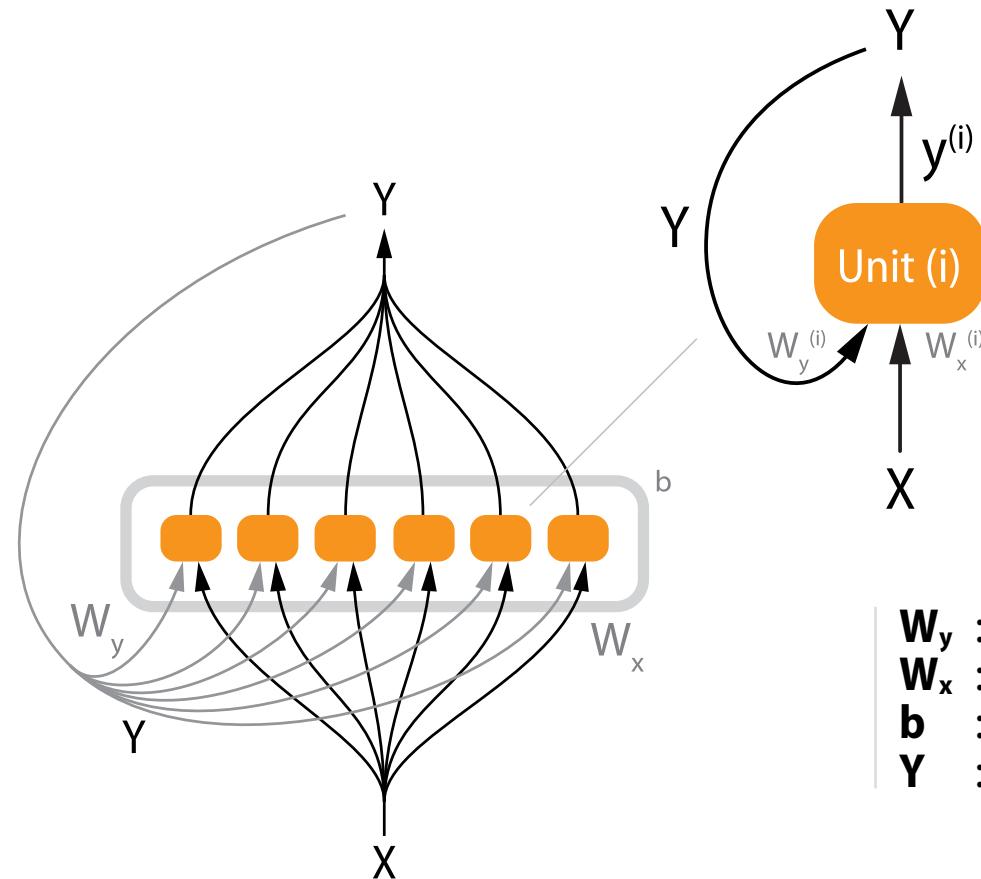
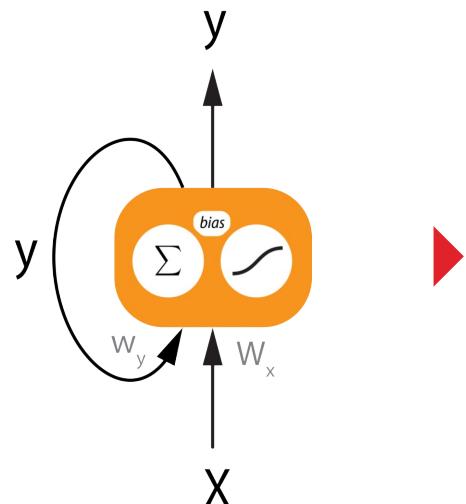


**$w_y$**  : is a scalar  
 **$W_x$**  : is a vector  
 **$b$**  : is a scalar  
 **$y$**  : is a scalar

$$y_{(t)} = \sigma(W_x^T \cdot X_{(t)} + w_y \cdot y_{(t-1)} + b)$$

# Recurrent Layer / Cell

## Recurrent Layer / Cell



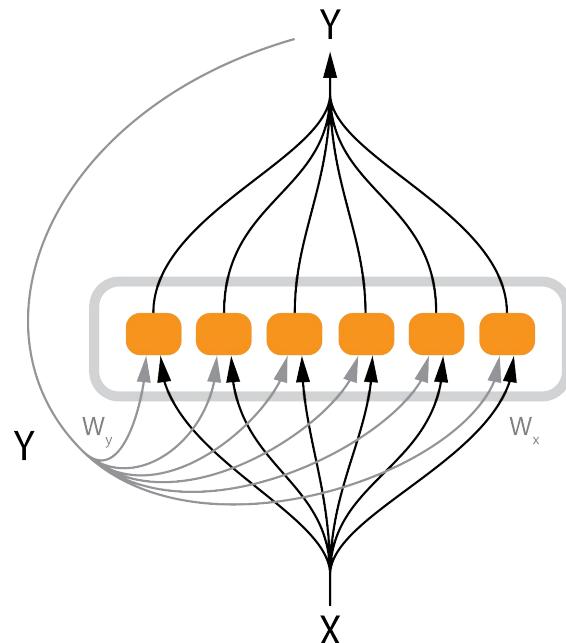
$W_y$  : is a tensor  
 $W_x$  : is a tensor  
 $b$  : is a vector  
 $Y$  : is a vector



Recurrent **neuron**.

Recurrent **Layer / Cell**

# Recurrent Layer / Cell



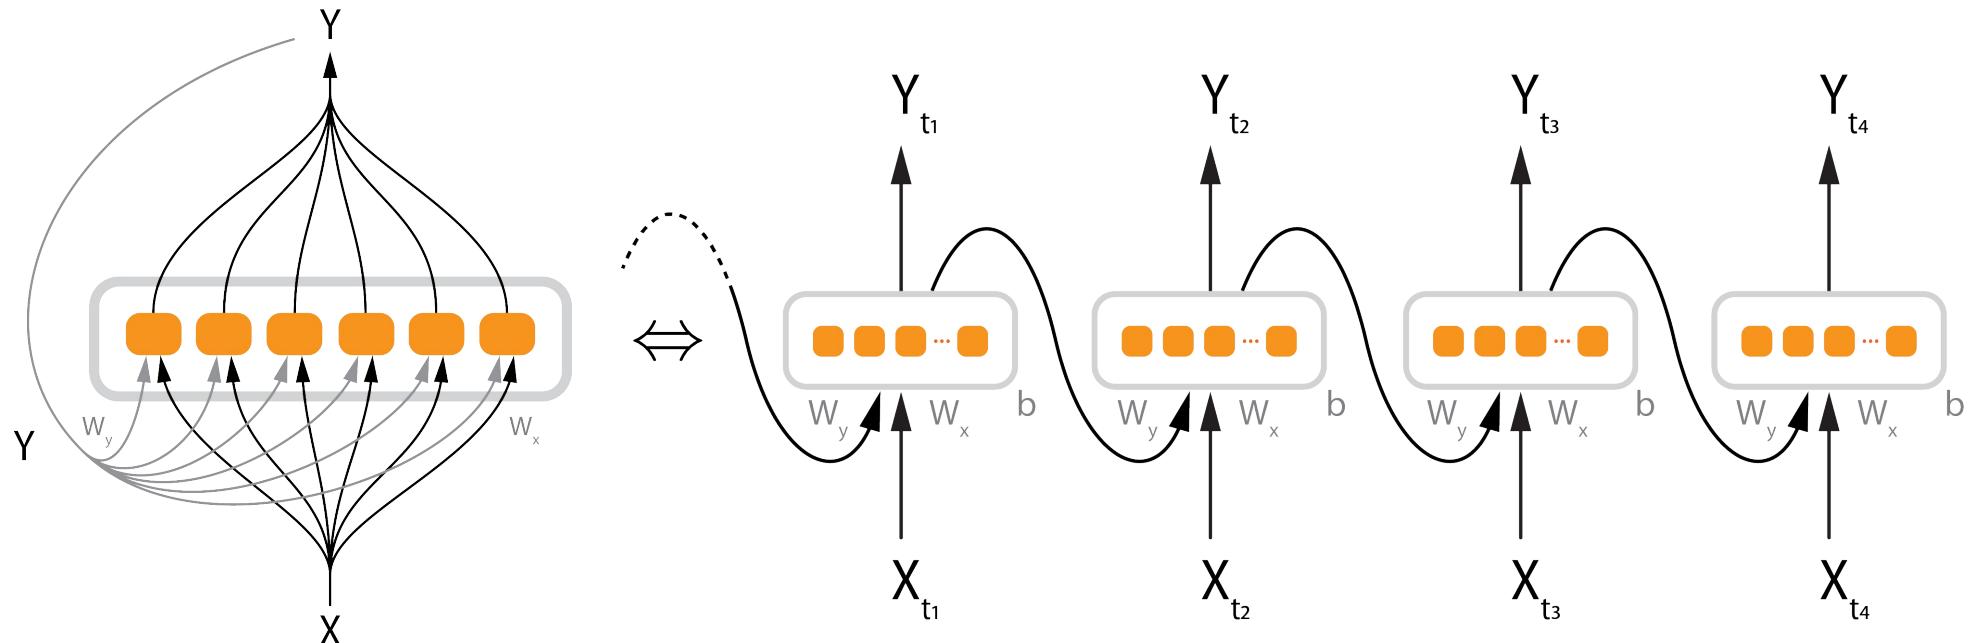
$$Y_{(t)} = \phi(W_x^T \cdot X_{(t)} + W_y^T \cdot Y_{(t-1)} + b)$$

**$W_y$**  : is a tensor  
 **$W_x$**  : is a tensor  
 **$b$**  : is a vector  
 **$Y$**  : is a vector

# Recurrent Layer / Cell

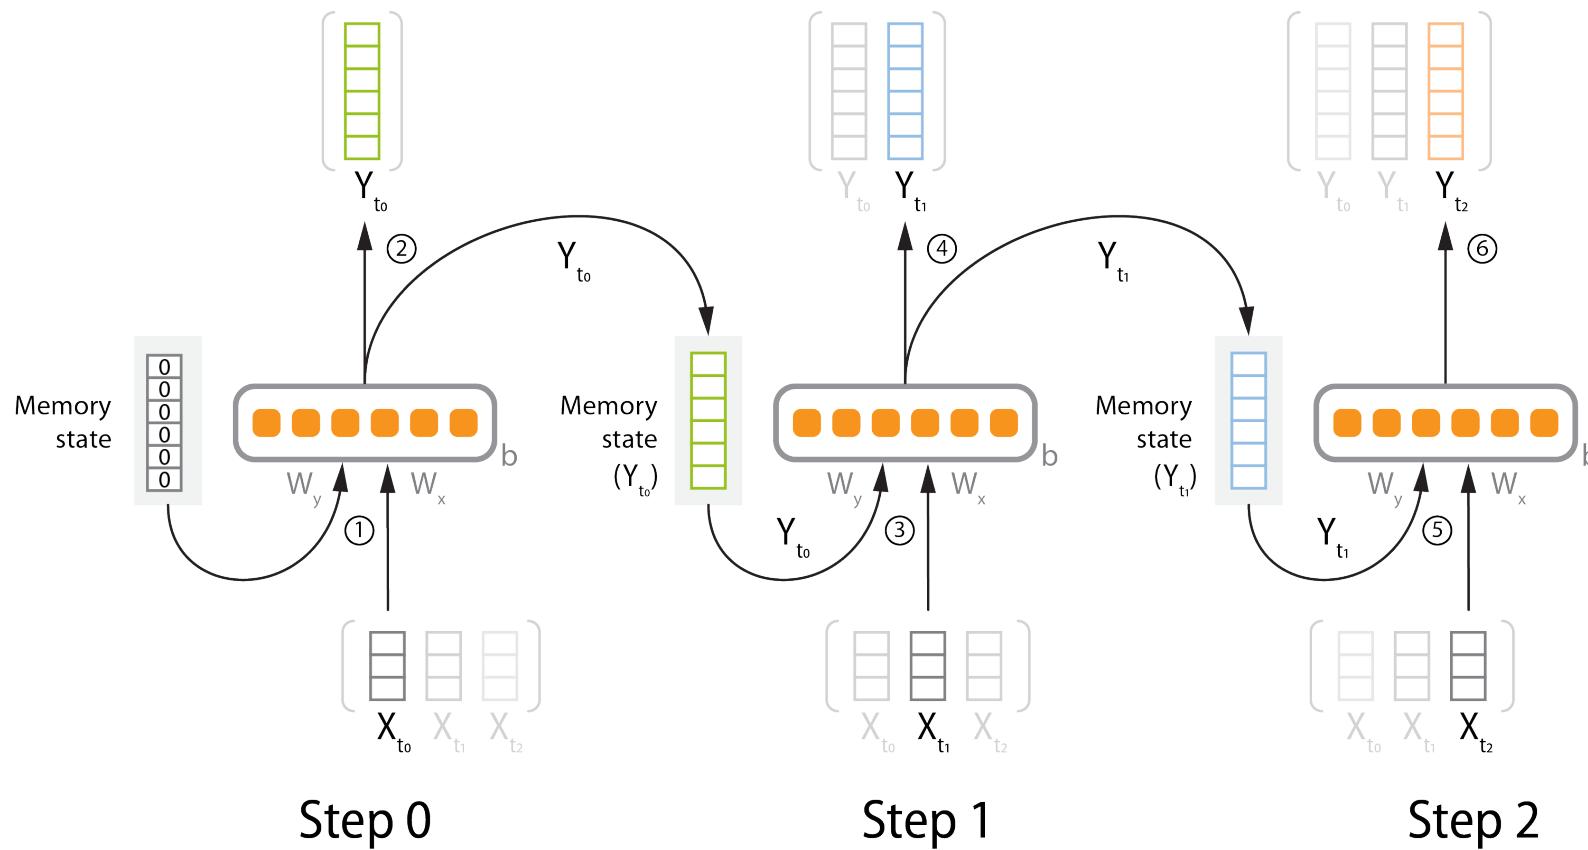
**$W_y$**  : is a tensor  
 **$W_x$**  : is a tensor  
 **$b$**  : is a vector  
 **$Y$**  : is a vector

## Recurrent Layer / Cell



$$Y_{(t)} = \phi(W_x^T \cdot X_{(t)} + W_y^T \cdot Y_{(t-1)} + b)$$

# Recurrent Layer / Cell



Shapes :

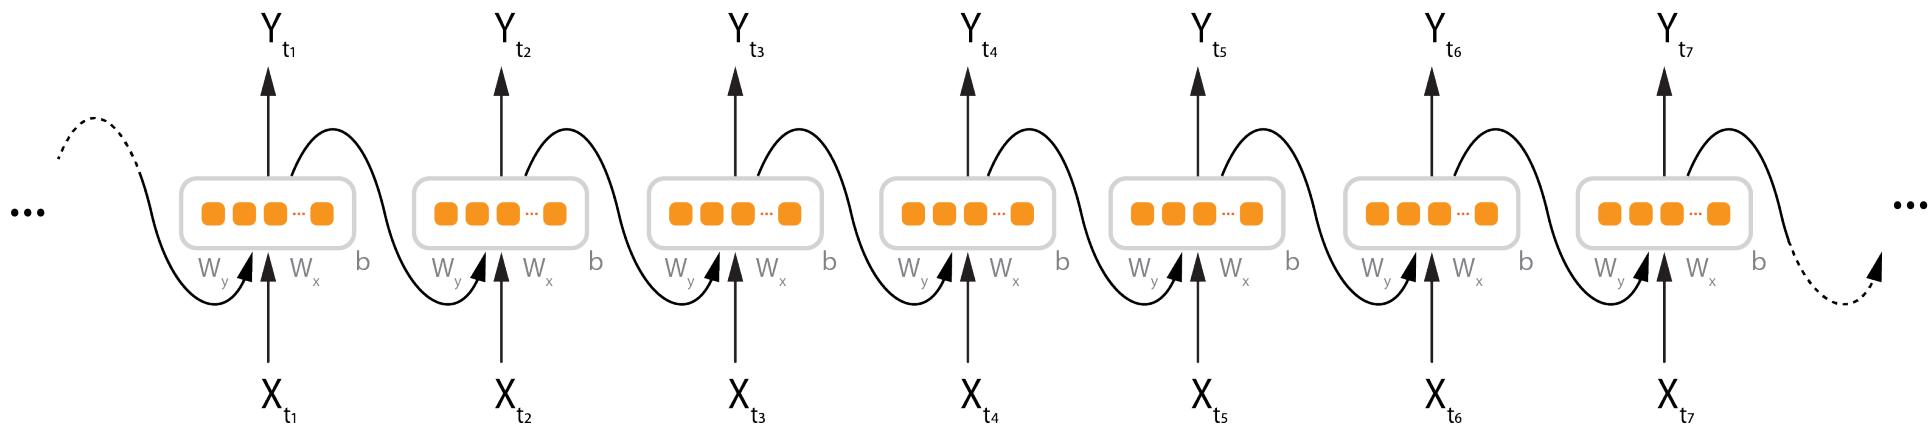
$W_x : (\text{nb}_{\text{units}}, \text{X size})$

$W_y : (\text{nb}_{\text{units}}, \text{nb}_{\text{units}})$

$Y : (\text{nb}_{\text{units}})$

$$Y_{(t)} = \Phi (W_x^T \cdot X_{(t)} + W_y^T \cdot Y_{(t-1)} + b)$$

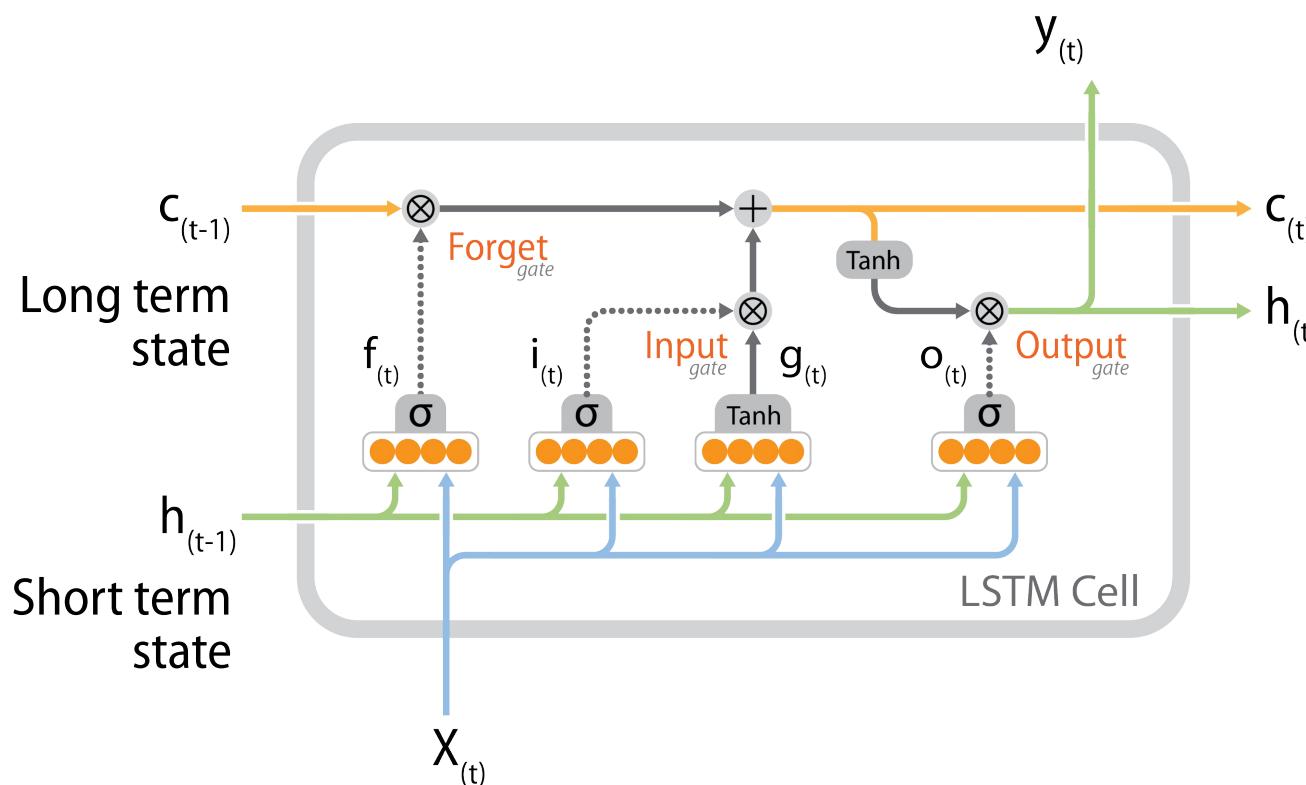
But....



Slow convergence,  
Short memory,  
Vanishing / exploding gradients

⇒ In short...  
it **doesn't**  
**work !**

# Long Short-Term Memory (LSTM)



$$\begin{aligned} f_{(t)} &= \sigma(W_{xf}^T X_{(t)} + W_{hf}^T h_{(t-1)} + b_f) \\ i_{(t)} &= \sigma(W_{xi}^T X_{(t)} + W_{hi}^T h_{(t-1)} + b_i) \\ g_{(t)} &= \tanh(W_{xg}^T X_{(t)} + W_{hg}^T h_{(t-1)} + b_g) \\ o_{(t)} &= \sigma(W_{xo}^T X_{(t)} + W_{ho}^T h_{(t-1)} + b_o) \\ c_{(t)} &= f_{(t)} \otimes c_{(t-1)} + i_{(t)} \otimes g_{(t)} \\ y_{(t)} &= h_{(t)} = o_{(t)} \otimes \tanh(c_{(t)}) \end{aligned}$$

with :

|                    |                    |                                 |
|--------------------|--------------------|---------------------------------|
| $X_{(t)}$          | $\in \mathbb{R}^d$ | input vector                    |
| $f_{(t)}$          | $\in \mathbb{R}^h$ | forget gate's activation vector |
| $i_{(t)}$          | $\in \mathbb{R}^h$ | input gate's activation vector  |
| $o_{(t)}$          | $\in \mathbb{R}^h$ | output gate's activation vector |
| $g_{(t)}$          | $\in \mathbb{R}^h$ | current entry vector            |
| $h_{(t)}, y_{(t)}$ | $\in \mathbb{R}^h$ | hidden state or output vector   |
| $c_{(t)}$          | $\in \mathbb{R}^h$ | cell state vector               |
| $\otimes$          |                    | Hadamard product                |
| $\sigma$           |                    | sigmoid function                |
| $W_k$              |                    | weights matrix                  |
| $b_k$              |                    | bias vector                     |

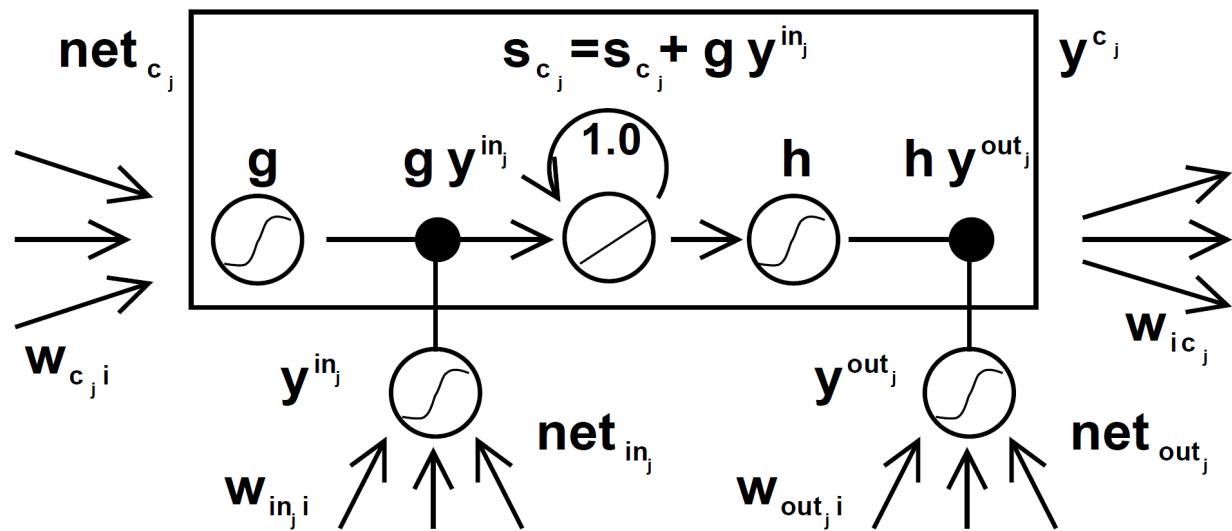
## Long Short-Term Memory

Sepp Hochreiter, Jürgen Schmidhuber, 1998  
<https://doi.org/10.1162/neco.1997.9.8.1735>

## Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation

Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, Yoshua Bengio, 2014  
<https://doi.org/10.48550/arXiv.1406.1078>

# Long Short-Term Memory (LSTM)



$$\begin{aligned} f_{(t)} &= \sigma(W_{xf}^T X_{(t)} + W_{hf}^T h_{(t-1)} + b_f) \\ i_{(t)} &= \sigma(W_{xi}^T X_{(t)} + W_{hi}^T h_{(t-1)} + b_i) \\ g_{(t)} &= \tanh(W_{xg}^T X_{(t)} + W_{hg}^T h_{(t-1)} + b_g) \\ o_{(t)} &= \sigma(W_{xo}^T X_{(t)} + W_{ho}^T h_{(t-1)} + b_o) \\ c_{(t)} &= f_{(t)} \otimes c_{(t-1)} + i_{(t)} \otimes g_{(t)} \\ y_{(t)} &= h_{(t)} = o_{(t)} \otimes \tanh(c_{(t)}) \end{aligned}$$

with :

|                                     |                                 |
|-------------------------------------|---------------------------------|
| $X_{(t)} \in \mathbb{R}^d$          | input vector                    |
| $f_{(t)} \in \mathbb{R}^h$          | forget gate's activation vector |
| $i_{(t)} \in \mathbb{R}^h$          | input gate's activation vector  |
| $o_{(t)} \in \mathbb{R}^h$          | output gate's activation vector |
| $g_{(t)} \in \mathbb{R}^h$          | current entry vector            |
| $h_{(t)}, y_{(t)} \in \mathbb{R}^h$ | hidden state or output vector   |
| $c_{(t)} \in \mathbb{R}^h$          | cell state vector               |
| $\otimes$                           | Hadamard product                |
| $\sigma$                            | sigmoid function                |
| $W_k$                               | weights matrix                  |
| $b_k$                               | bias vector                     |

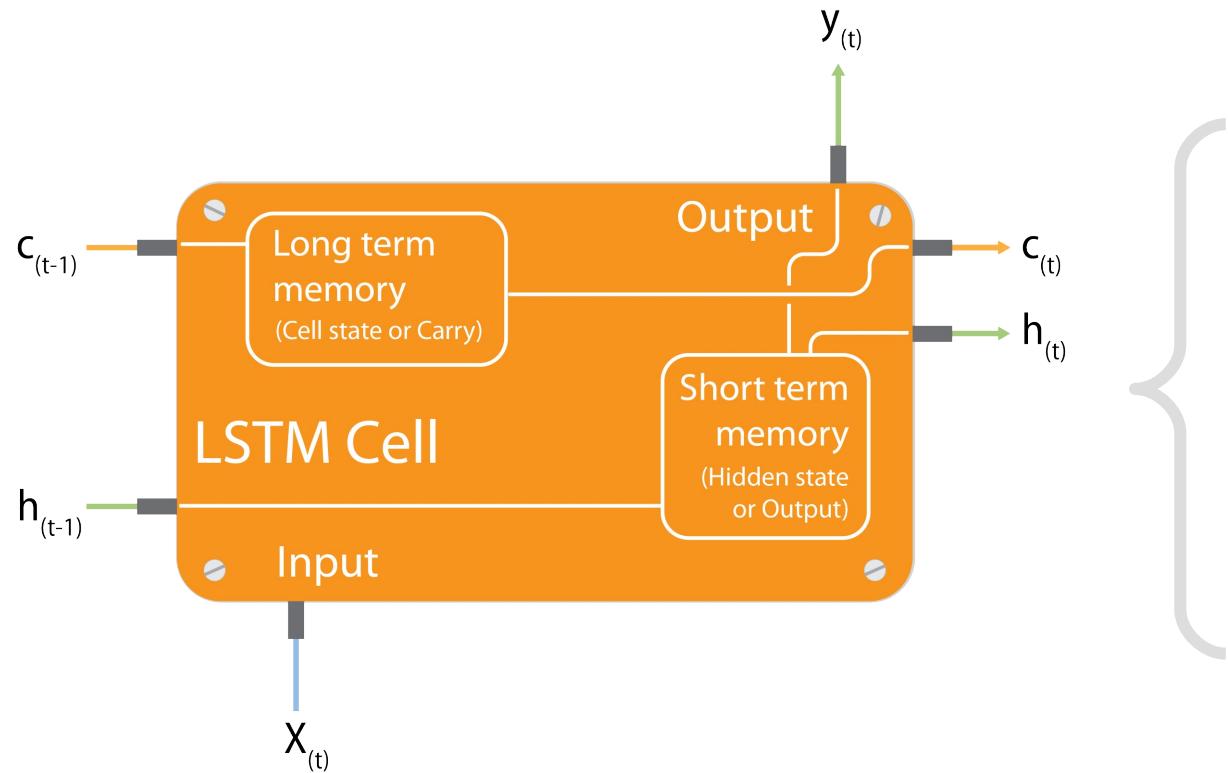
## Long Short-Term Memory

Sepp Hochreiter, Jürgen Schmidhuber, 1998  
<https://doi.org/10.1162/neco.1997.9.8.1735>

## Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation

Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, Yoshua Bengio, 2014  
<https://doi.org/10.48550/arXiv.1406.1078>

# Long Short-Term Memory (LSTM)



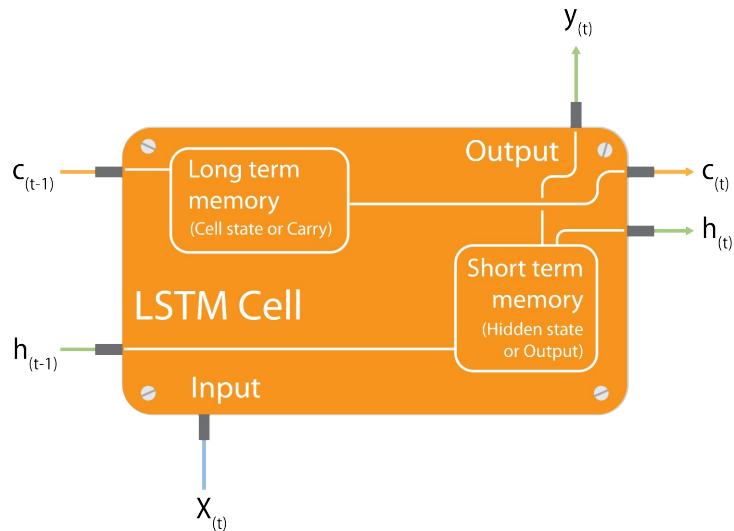
Keras  
PyTorch



**Long Short-Term Memory**  
Sepp Hochreiter, Jürgen Schmidhuber, 1998  
<https://doi.org/10.1162/neco.1997.9.8.1735>

**Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation**  
Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, Yoshua Bengio, 2014  
<https://doi.org/10.48550/arXiv.1406.1078>

# Long Short-Term Memory (LSTM)



## Serie to vector

```
lstm = keras.layers.LSTM(18, return_sequences=True, return_state=True)
output, memory_state, carry_state = lstm(input)
```

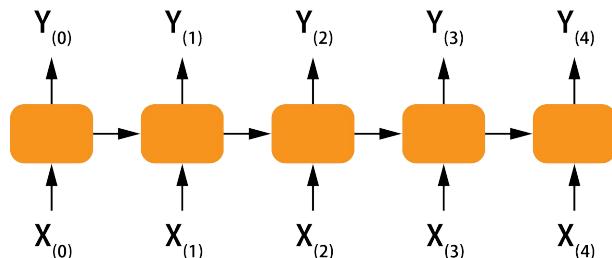
Inputs shape is : (32, 20, 8)  
Output shape is : (32, 18)

Output shape : (32, 20, 18)  
Memory state : (32, 18)  
Carry state : (32, 18)

More about :

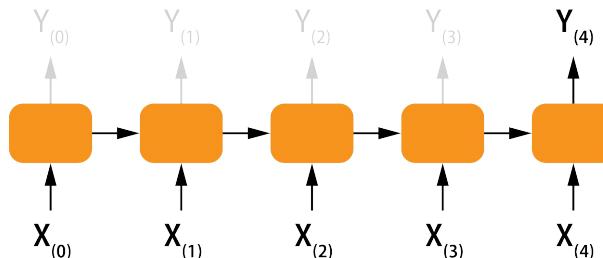
[https://keras.io/api/layers/recurrent\\_layers/lstm/](https://keras.io/api/layers/recurrent_layers/lstm/)

# Reccurent Neural Network (RNN)



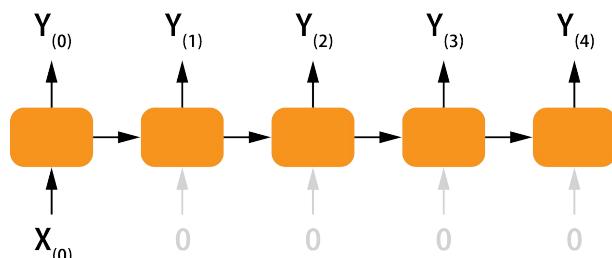
**Serie → serie**

Example : Time serie prediction



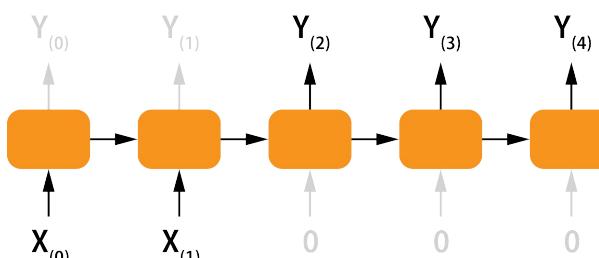
**Serie → vector**

Example : Sentiment analysis



**Vector → serie**

Example : Image annotation



**Encoder → decoder**

Example : Language Translation

# How to predict a sequence ?

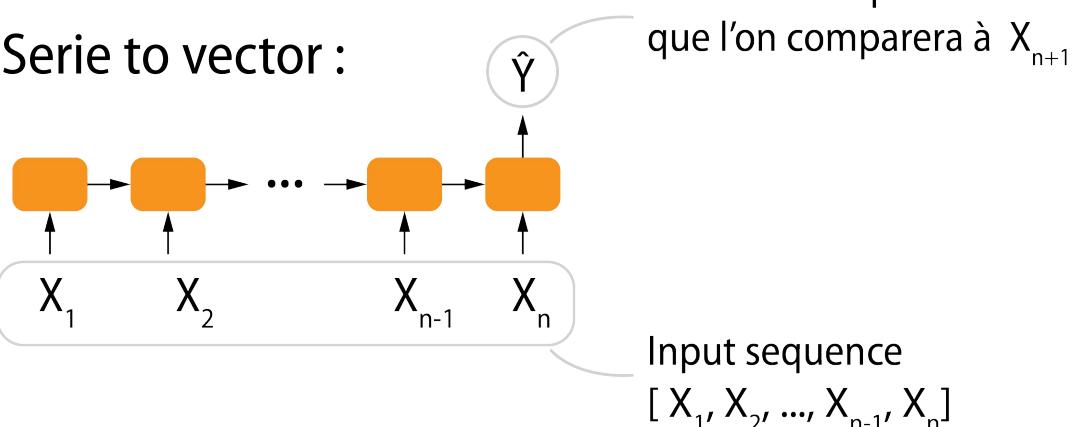
Known sequence :

[ $X_1, X_2, X_3, X_4, \dots, X_n, X_{n+1}$ ]

Input Sequence :  $X = [X_1, X_2, \dots, X_n]$

Expected output :  $Y = [X_{n+1}]$

Serie to vector :



Objective :

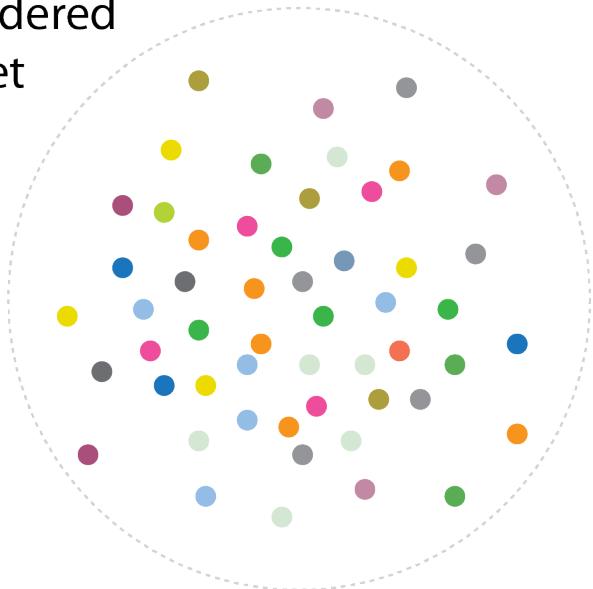
Train our RNN network to **predict** the **n+1** vector of our input sequence :

# Preparation of sequence data



The distribution of data between trains and test series must be comparable.

Not ordered  
Dataset



Shuffle

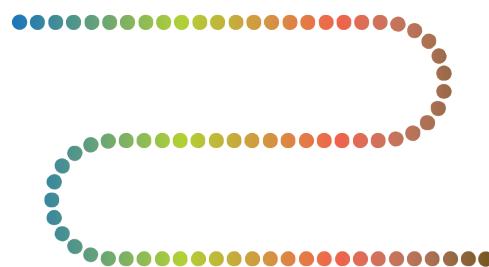


# Preparation of sequence data



Can the past  
explain the future ?

Ordered dataset

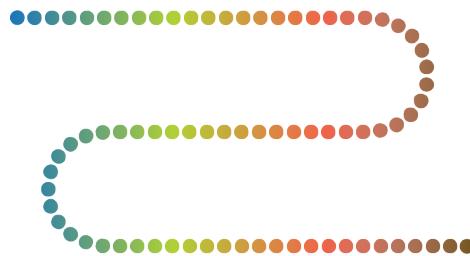


Train data



# Preparation of sequence data

Ordered dataset



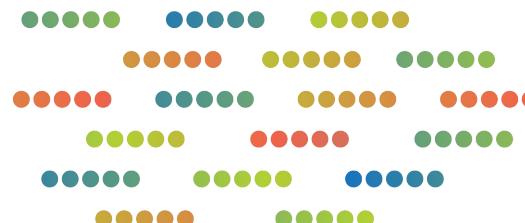
Train data



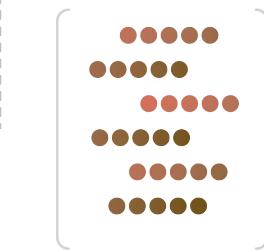
Past

Future

Test data



Train set

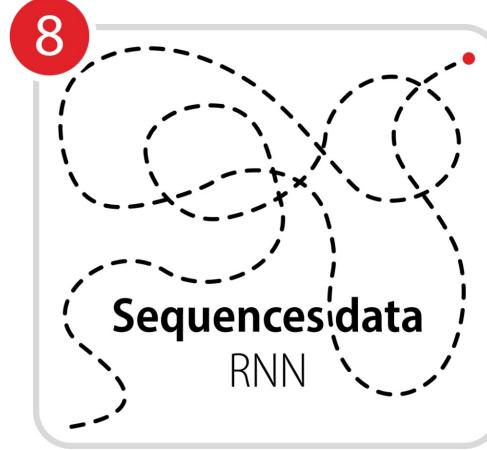


Test set



Note that a data generator should be very useful !



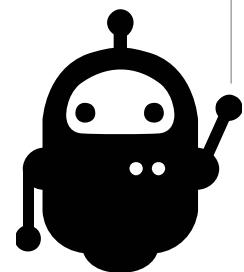


## Part 2

- 1 **Sequences data**  
Recurrent Neural Network  
LSTM and GRU

- 2 **Example 1 : LadyBug**  
→ Prediction of a virtual trajectory





# Time series with RNN

Notebook : [\[LADYB1\]](#)

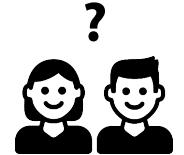
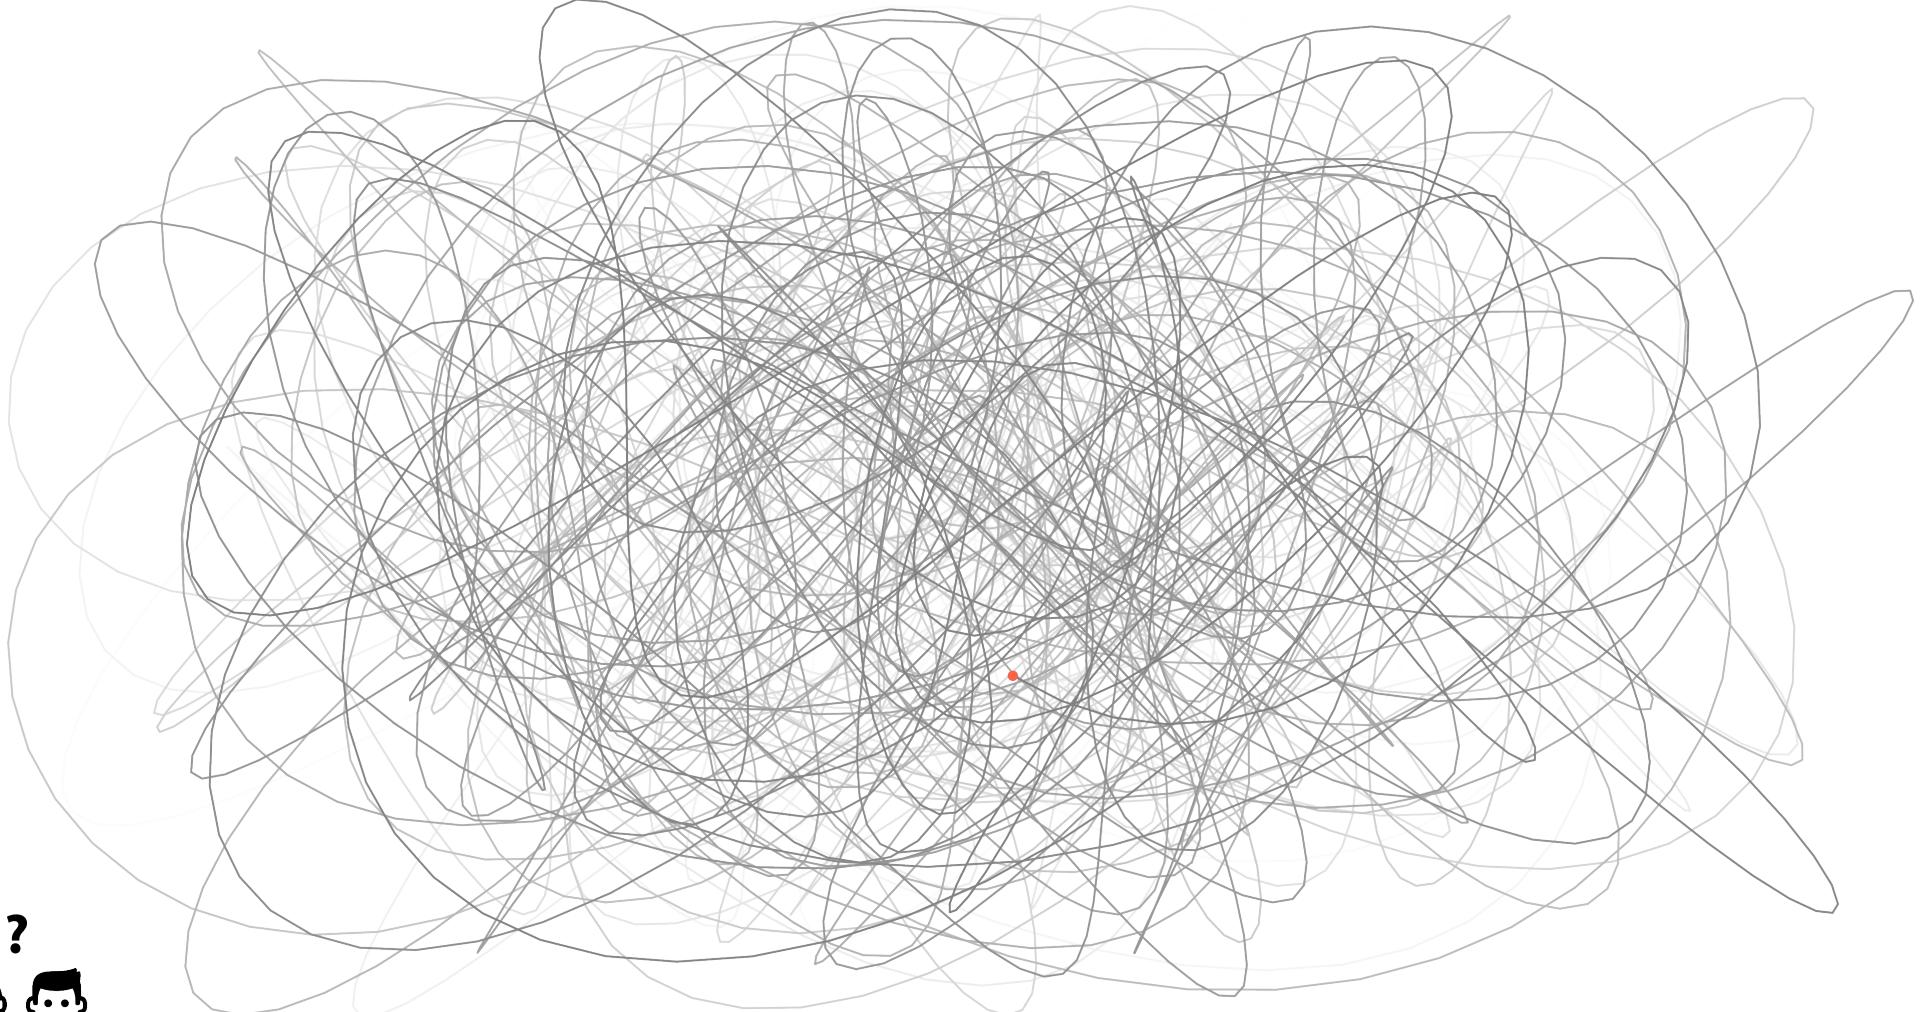
**Objective :**

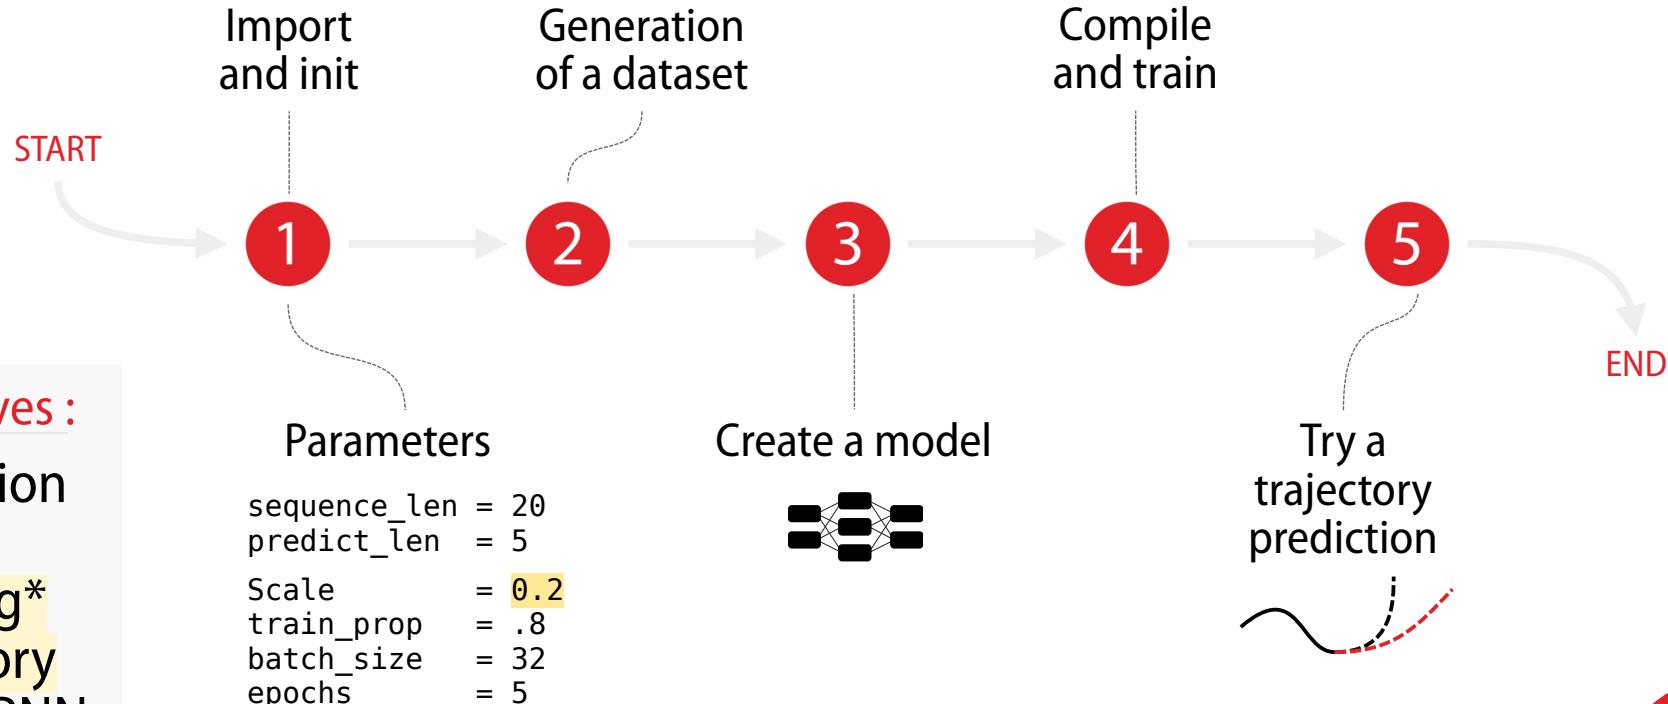
Prediction of a 2D ladybug trajectory with a RNN

**Dataset :**

Generated trajectory







\*Disclaimer : No ladybugs were harmed during the making of this notebook.





## Little things and concepts to **keep in mind**

- Can the **past explain the future** ?
- Understand the data, again and again !
- Beware of **overfitting**
- Remember that **Pandas** is good for you !
- The json files are cool, too
- **Preparing** your data can cost 70% of the work
- Think about **data generators**
- Matplotlib are also very good for you !
- There is a lot of **sequential data**
- Not everything can uses **GPUs...**

But the story isn't over...  
...and this was just the beginning!

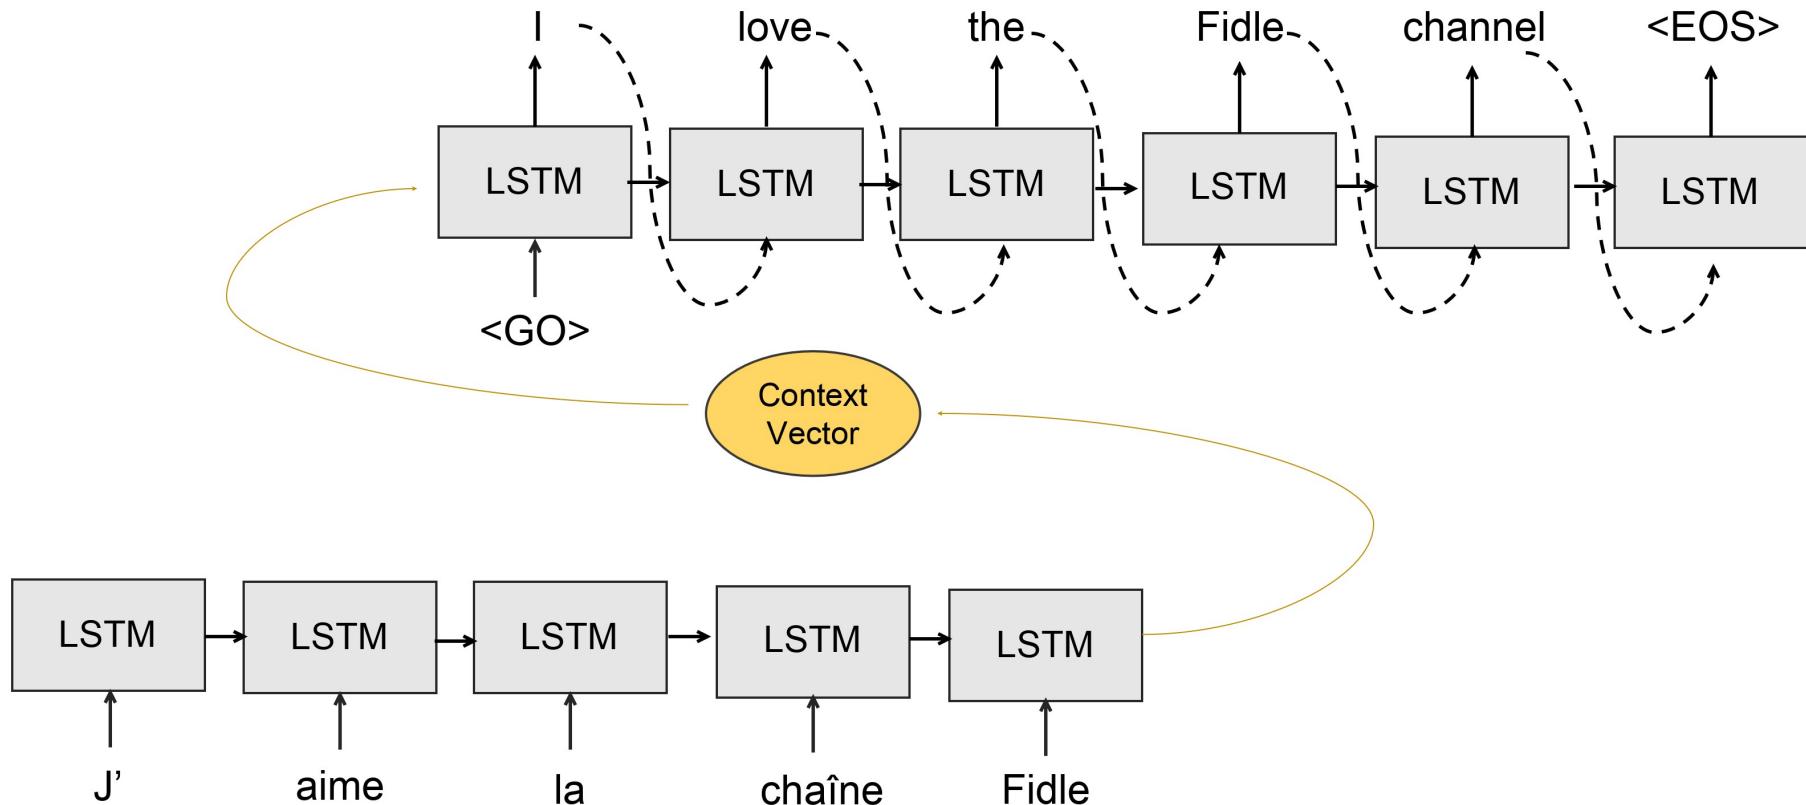


لَكِنَّ الْقَصْةَ لَمْ تُنْتَهِ بَعْدَ...  
وَهَذِهِ كَانَتْ لِبَدْلِيَّةً فَقْطَ...

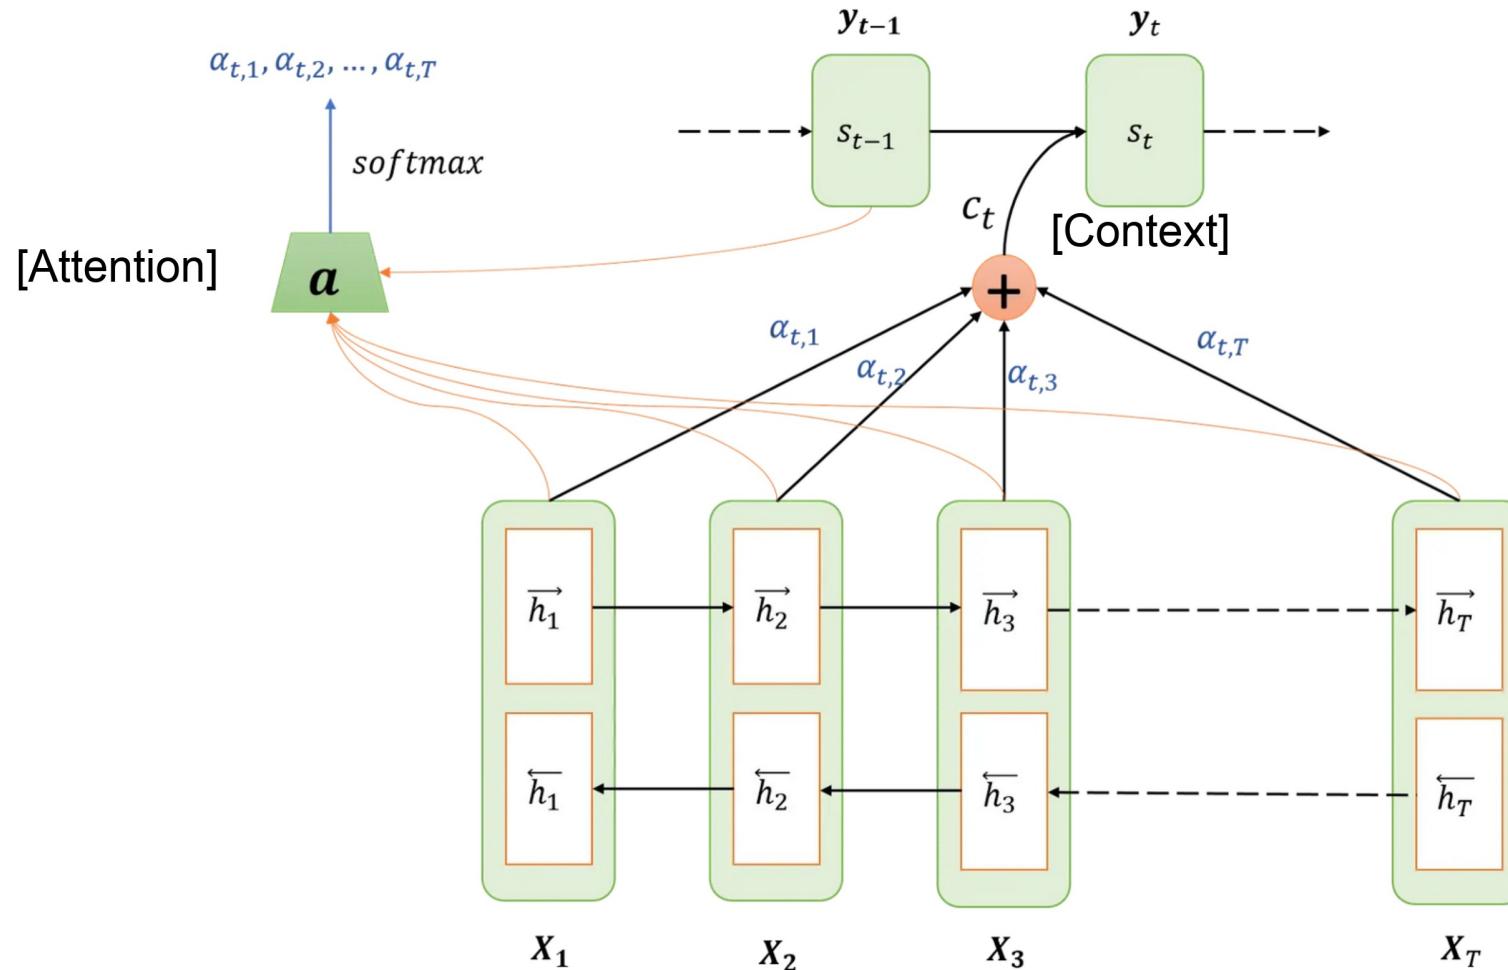
하지만 이야기는 끝나지 않았습니다 ...  
... 그리고 이것은 시작에 불과합니다 !



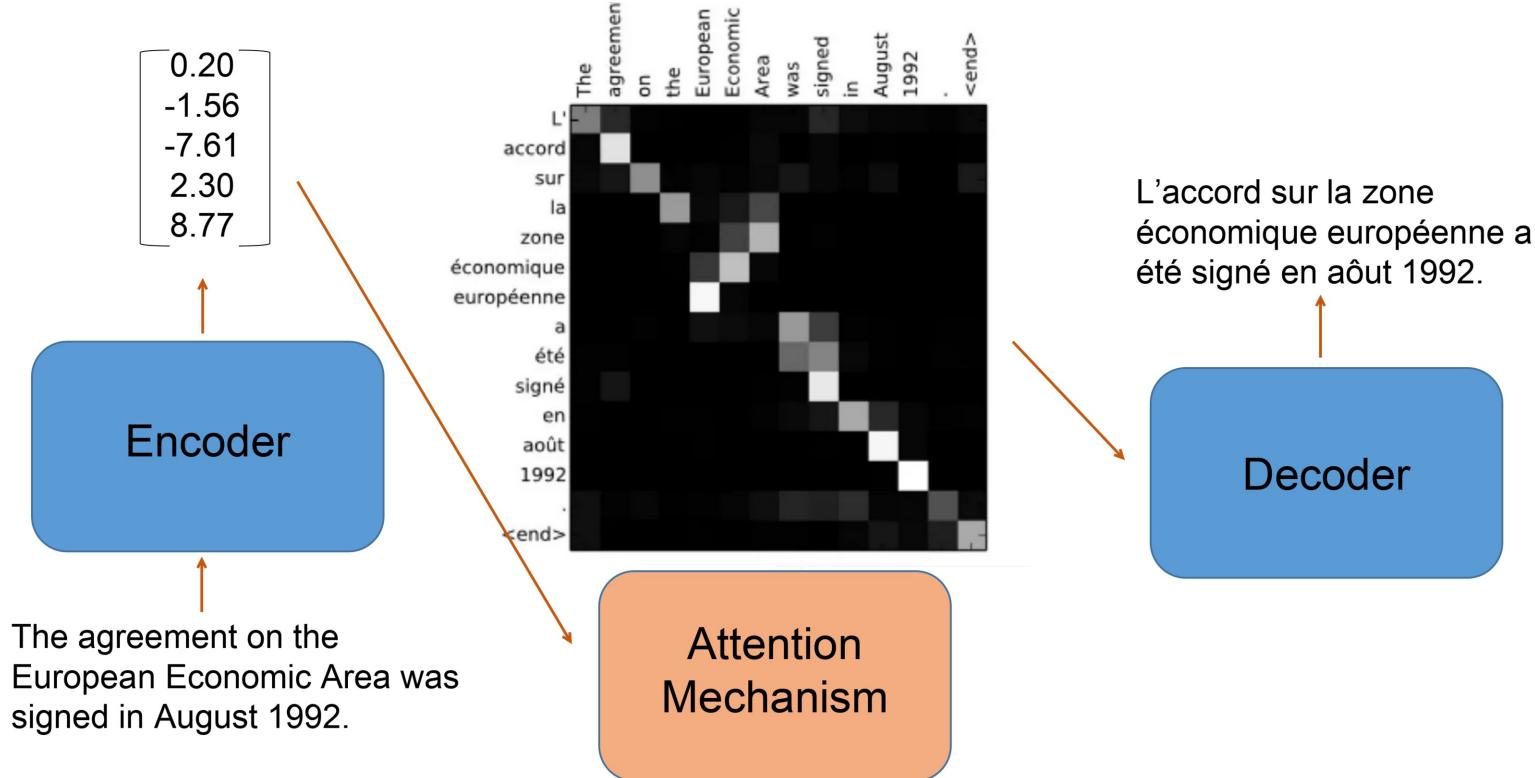
# Simple translator, sequence to sequence



# Attention Mechanism



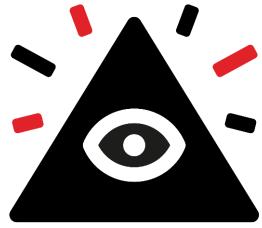
# Attention Mechanism



Next, on Fidle :



9



«Attention is  
All You Need»

Transformers



Jeudi  
08  
Février  
à 14h00